

컴퓨터 비전 초격차: 이론/실습

2 Preliminary – Linear algebra

Index

Preliminary

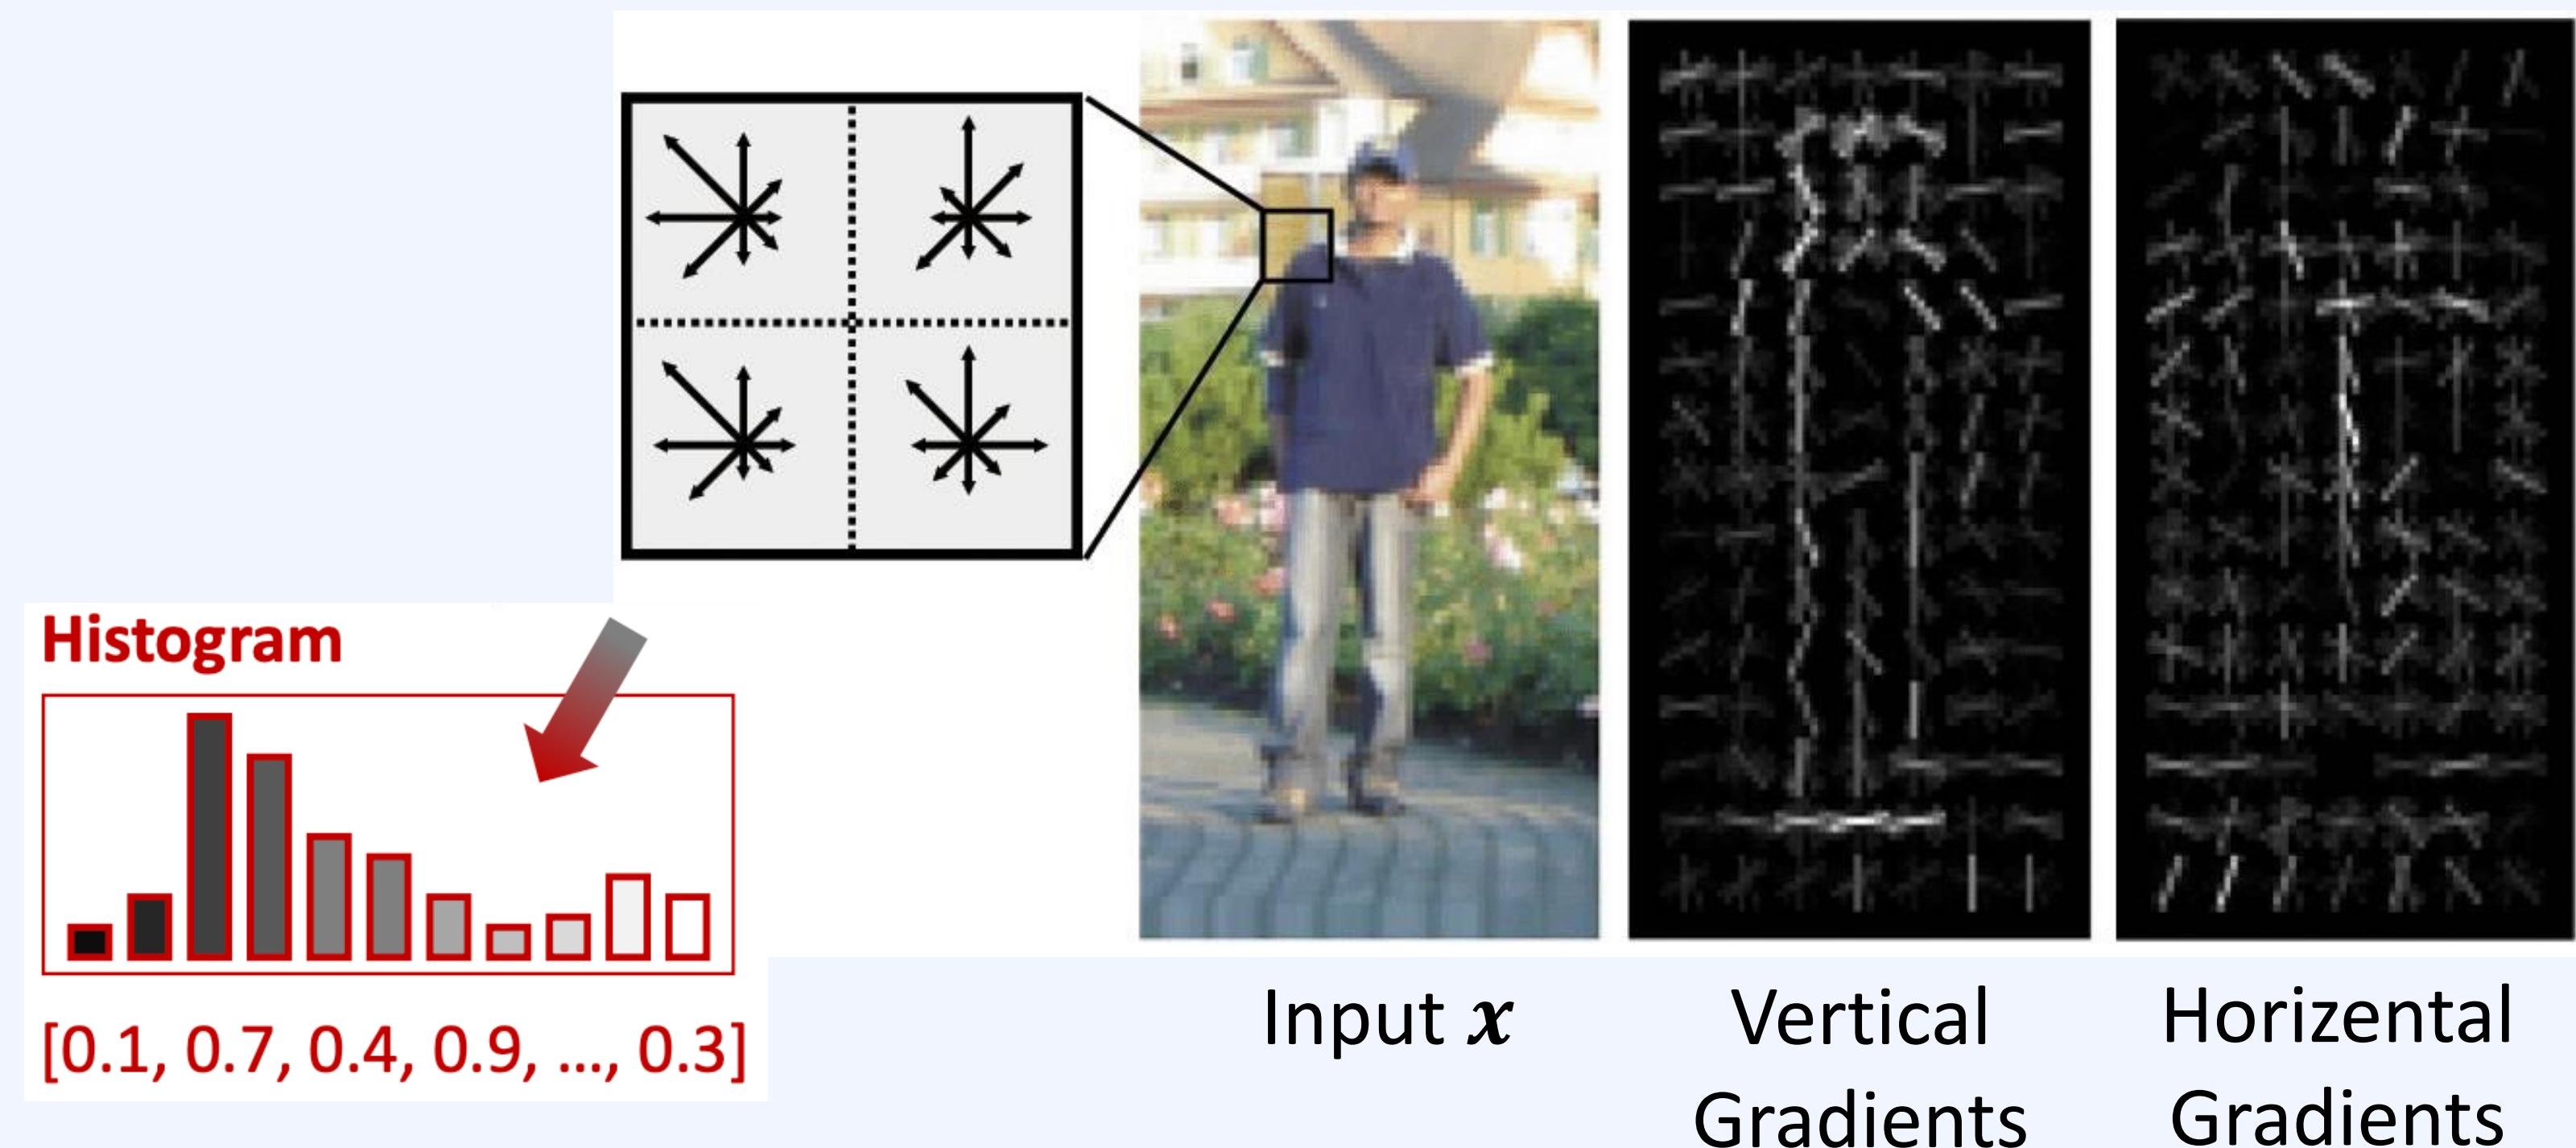
Linear Algebra
Basic

- Linear Algebra in Computer Vision
- Vector
 - Operations
 - Norm
 - Linear Dependency
 - Basis
- Matrix
 - Oprations
 - Rank
 - Determinant
 - Inverse Matrix
- Eigen Decomposition

Most of the slides are inspired by POSTECH
AIGS549 class (Prof. Suha Kwak)

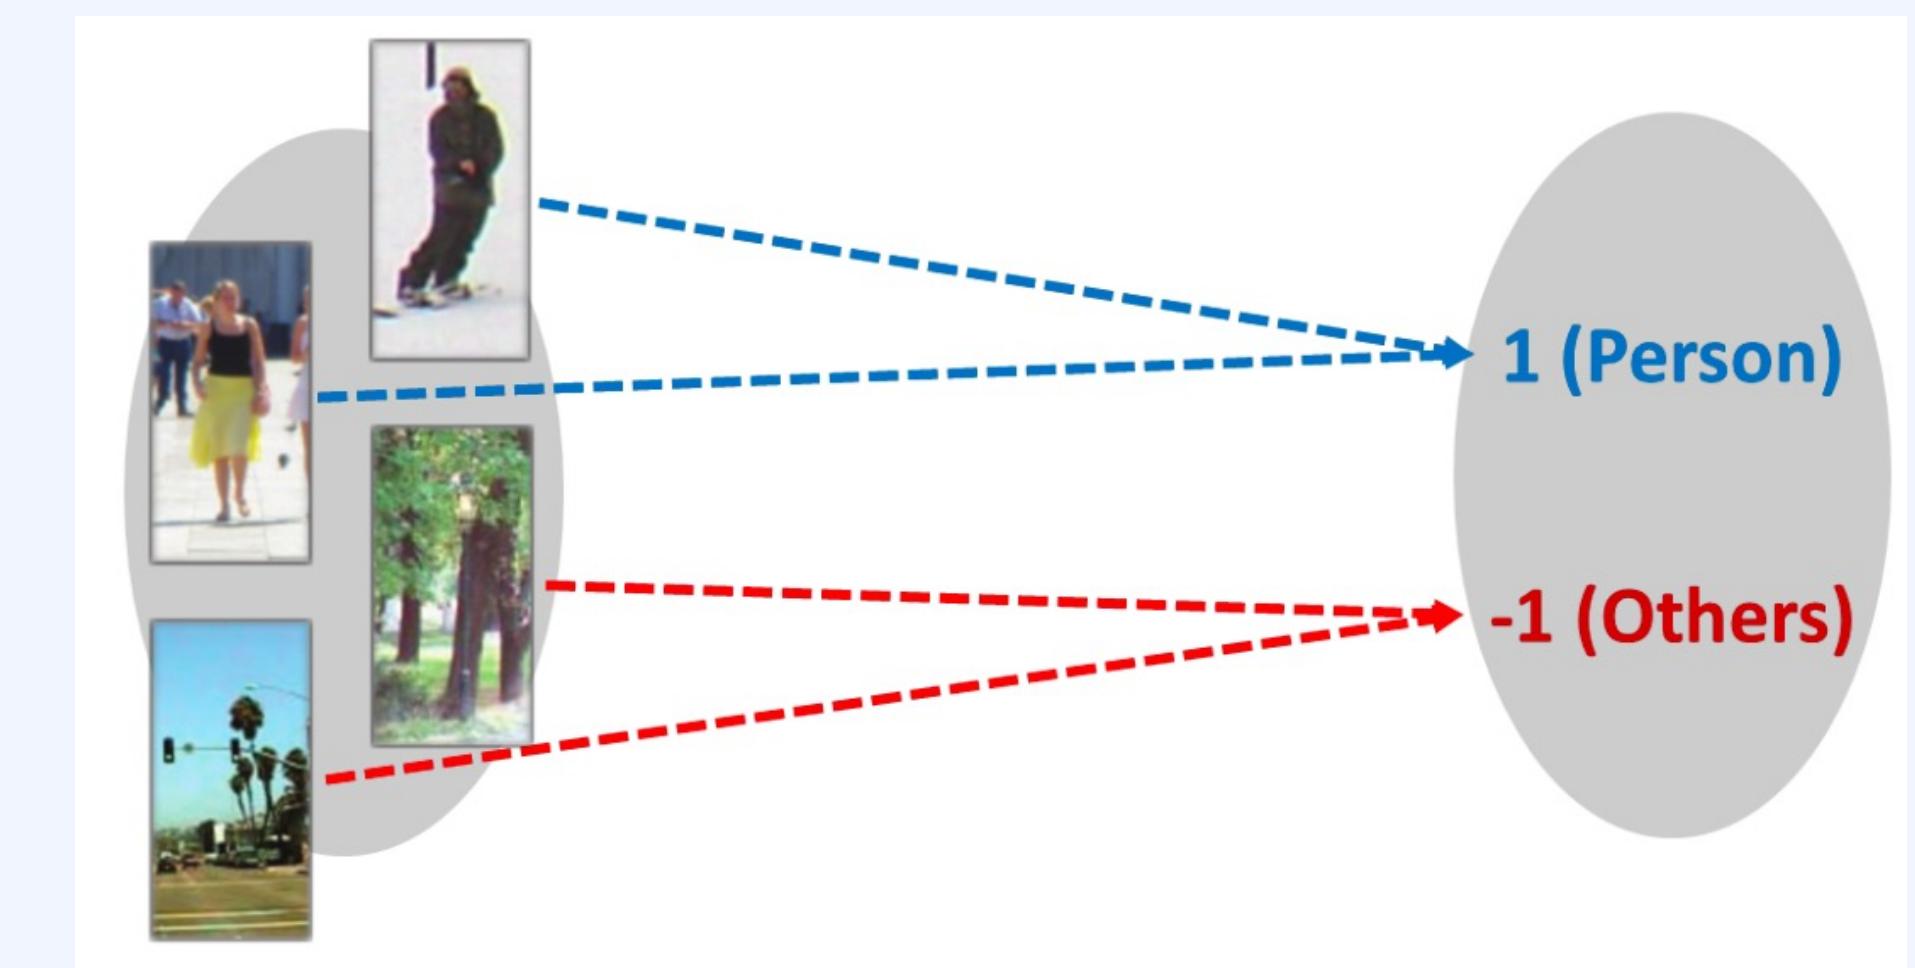
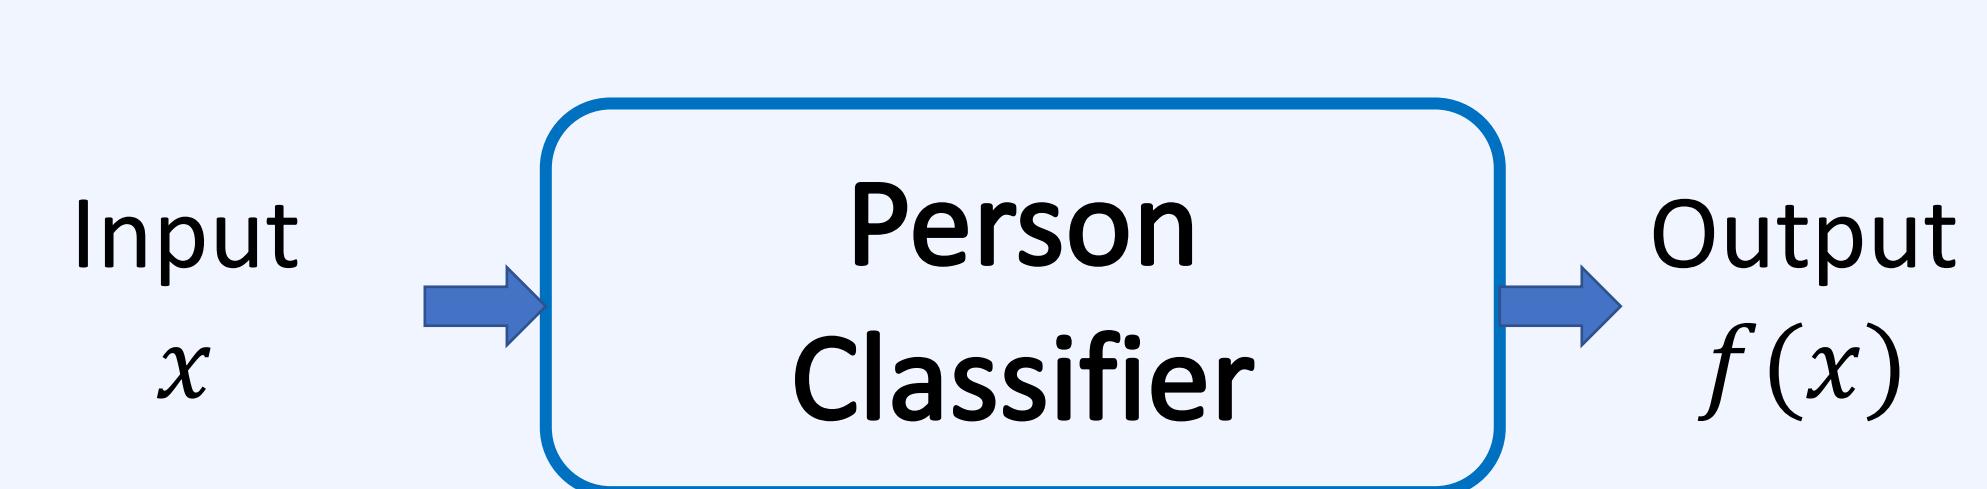
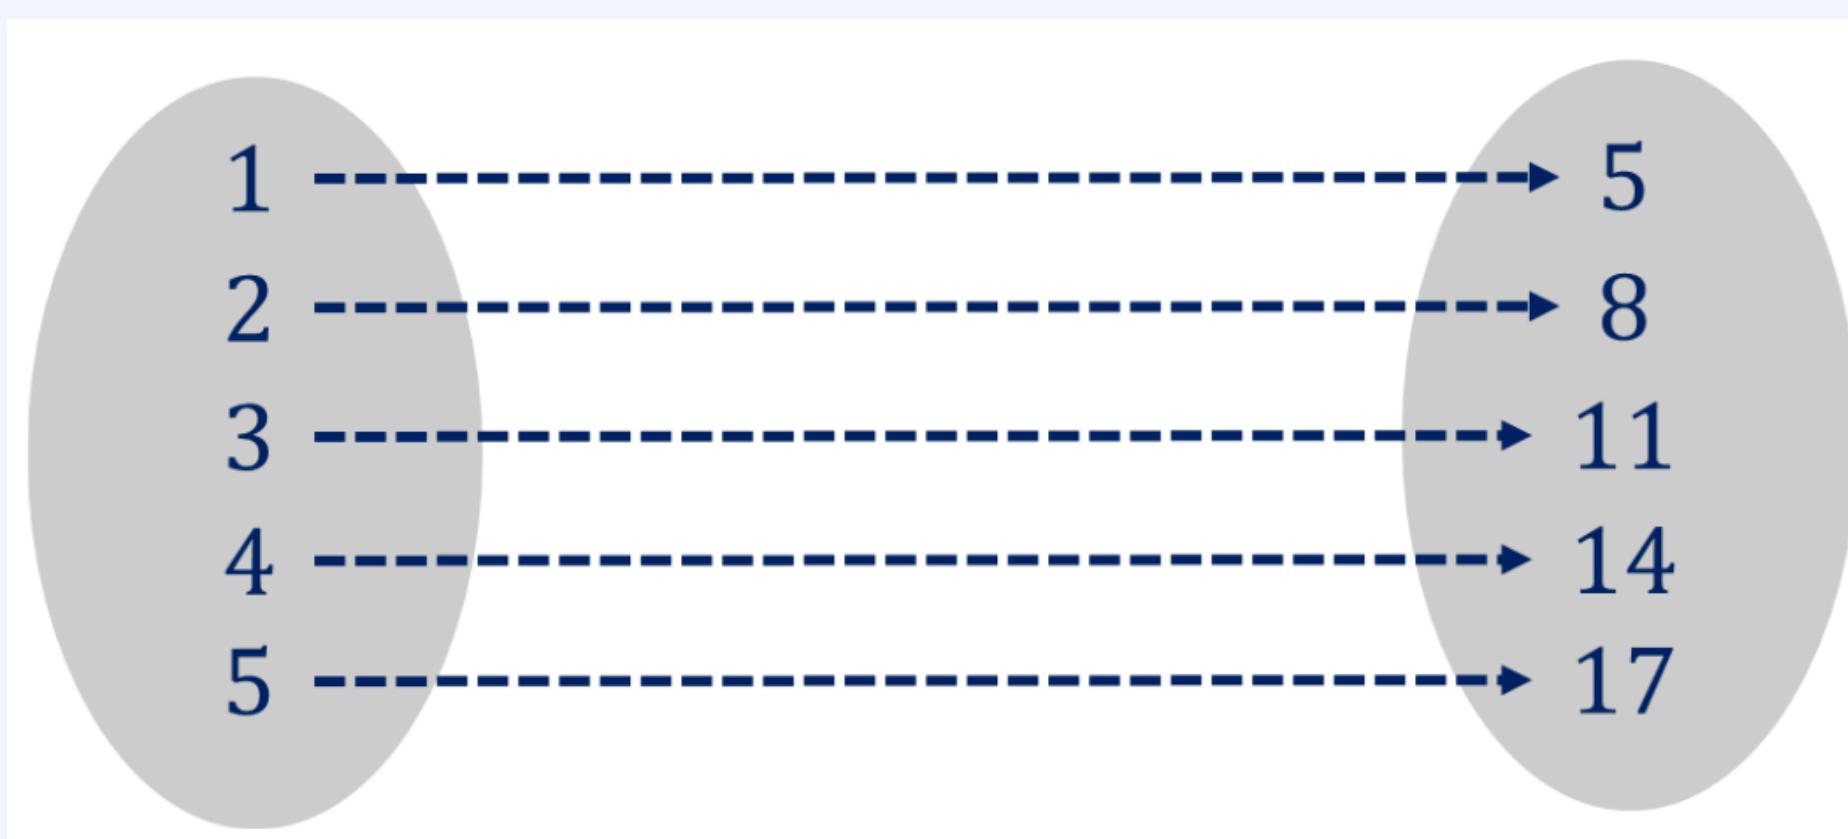
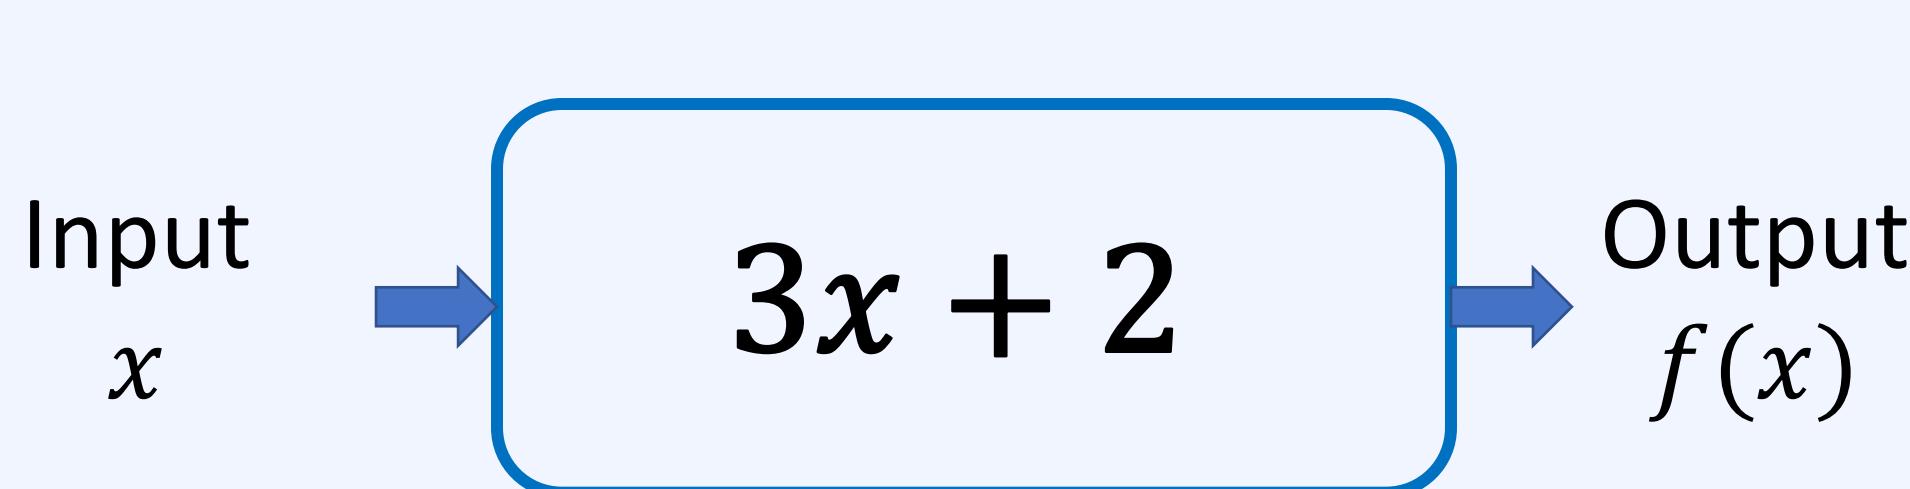
Linear Algebra in Computer Vision

- Mathematics in **vector space**
- E.g., Image understanding = **Image description + Decision Making (Vector)**



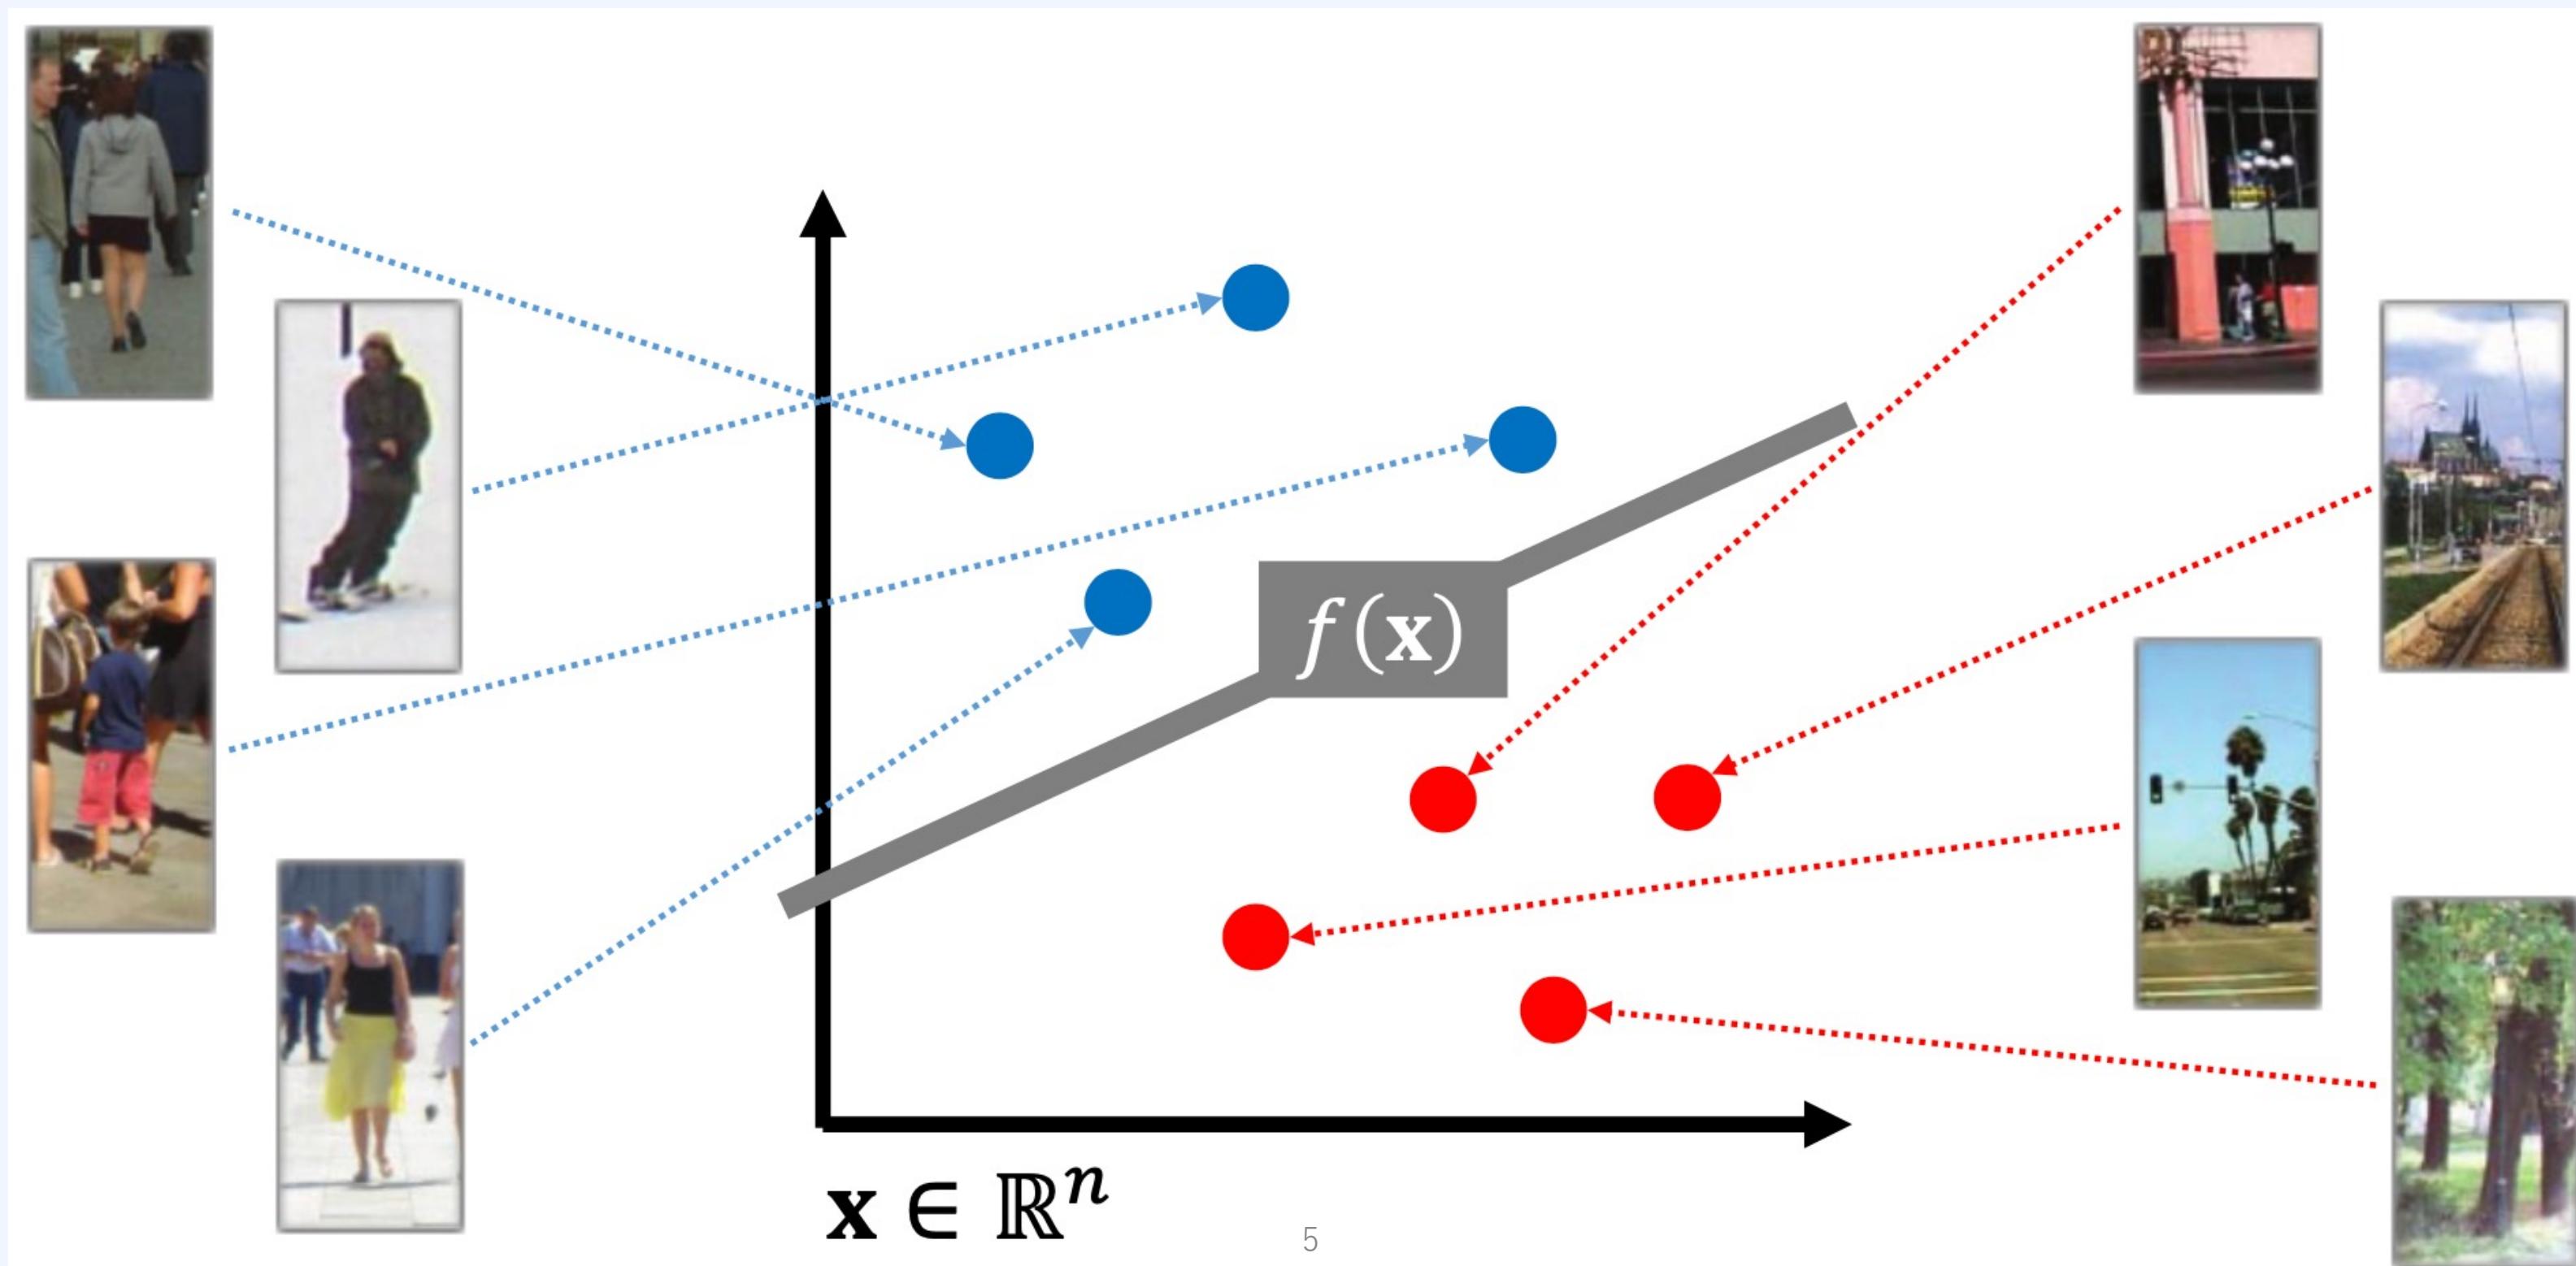
Linear Algebra in Computer Vision

- Mathematics in **vector space**
- E.g., Image understanding = **Image description (Vector)** + **Decision Making (Function)**



Linear Algebra in Computer Vision

- Mathematics in **vector space**
- E.g., Image understanding = **Image description (Vector)** + **Decision Making (Function)**



Linear Algebra in Computer Vision

- Mathematics in **vector space**
- Frequently used in
 - Image description
 - Computation of similarities and distances
 - Finding algebraic solutions
 - Transformations
 - Optimization

Vector

- Geometric object that has both a **magnitude** and **direction**
- Notation

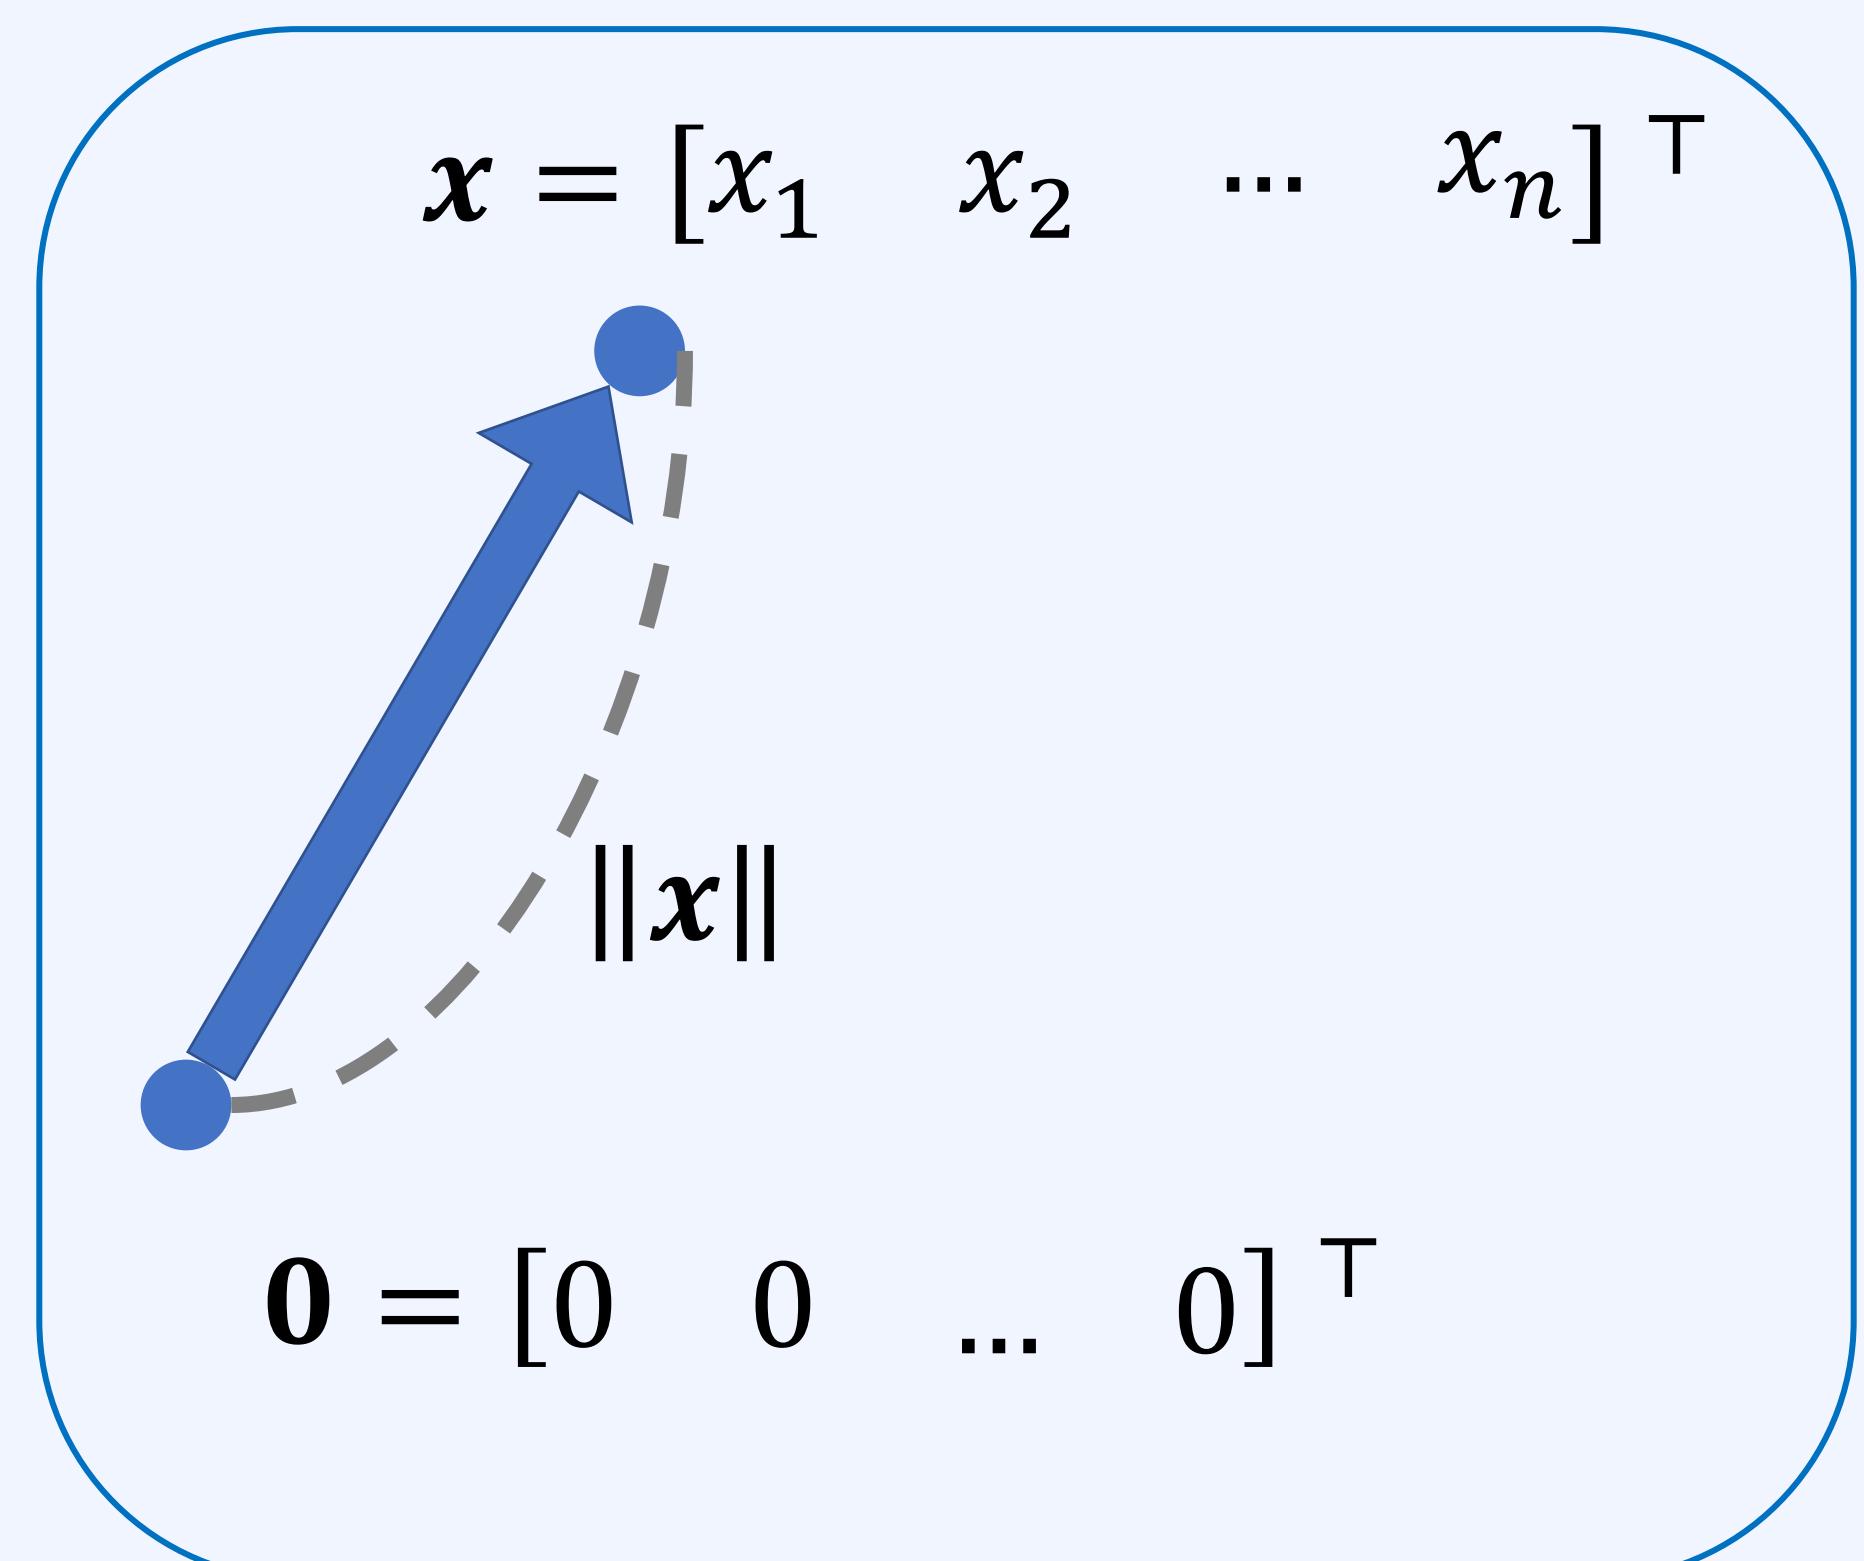
$$\boldsymbol{x} \in \mathbb{R}^n, \quad \boldsymbol{x} \in \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix}$$

- Transpose

$$\boldsymbol{x}^\top = [x_1 \ x_2 \ \dots \ x_n]$$

- Magnitude

$$\|\boldsymbol{x}\| = \sqrt{x_1^2 + x_2^2 + \dots + x_n^2}$$



Vector Operations

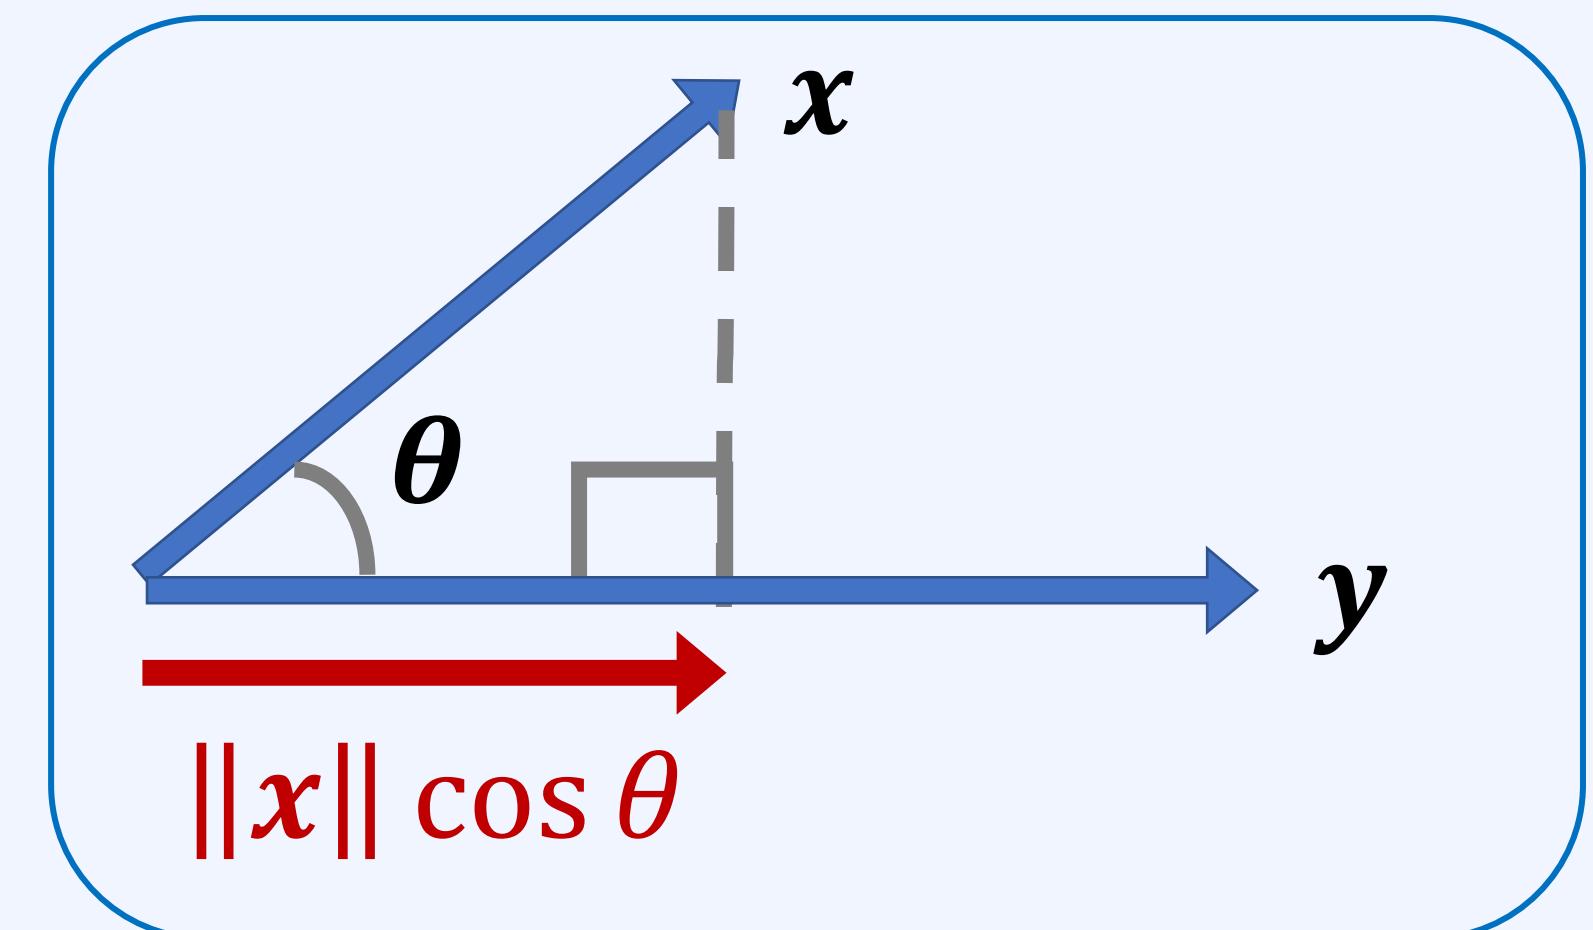
- Inner product (a.k.a., dot product)

$$\mathbf{x} = [x_1 \quad x_2 \quad \dots \quad x_n]^T \quad \mathbf{y} = [y_1 \quad y_2 \quad \dots \quad y_n]^T$$

$$\begin{aligned}\mathbf{x} \cdot \mathbf{y} &= \|\mathbf{x}\| \|\mathbf{y}\| \cos \theta \\ &= \sum_{i=1}^n x_i y_i = x_1 y_1 + \dots + x_n y_n = \mathbf{x}^T \mathbf{y}\end{aligned}$$

- Properties of inner product
 - Symmetry: $\mathbf{x} \cdot \mathbf{y} = \mathbf{y} \cdot \mathbf{x}$
 - Distributiveness: $(\mathbf{x} + \mathbf{y}) \cdot \mathbf{z} = \mathbf{x} \cdot \mathbf{z} + \mathbf{y} \cdot \mathbf{z}$
 - Linearity: $(\lambda \mathbf{x}) \cdot \mathbf{y} = \mathbf{x} \cdot (\lambda \mathbf{y}) = \lambda(\mathbf{x} \cdot \mathbf{y})$
 - Positive definite: $\mathbf{x} \cdot \mathbf{x} = \mathbf{x}^T \mathbf{x} > 0, \forall \mathbf{x} \neq \mathbf{0}$

$$\mathbf{x} \cdot \mathbf{x} = \mathbf{0} \Leftrightarrow \mathbf{x} = \mathbf{0}$$



Vector Operations

- Outer product

$$\mathbf{x} = [x_1 \quad x_2 \quad \dots \quad x_m]^\top \in \mathbb{R}^m$$

$$\mathbf{y} = [y_1 \quad y_2 \quad \dots \quad y_n]^\top \in \mathbb{R}^n$$

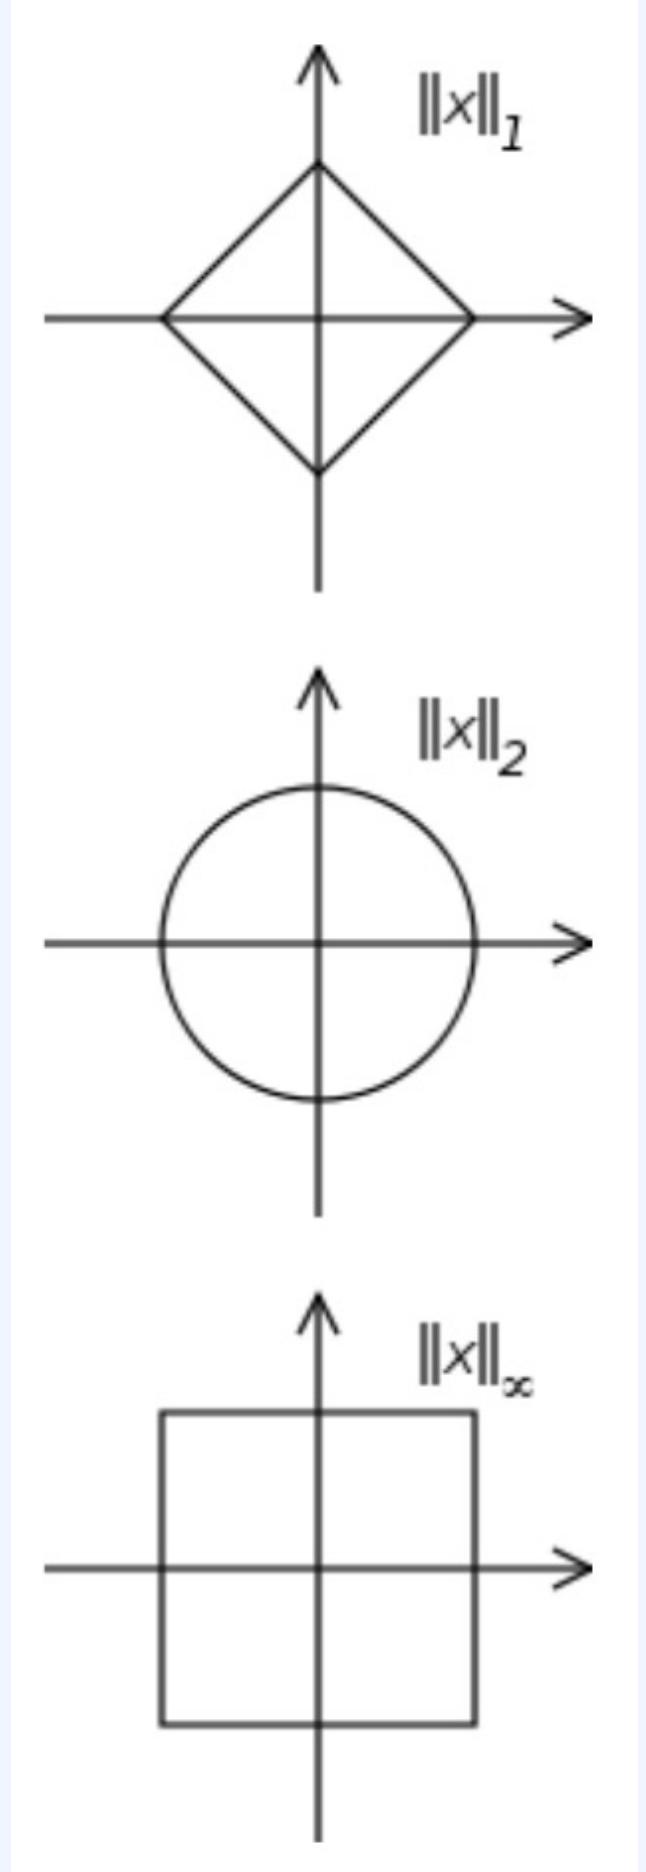
$$\mathbf{x} \otimes \mathbf{y} = \mathbf{x}\mathbf{y}^\top = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_m \end{bmatrix} [y_1 \quad y_2 \quad \dots \quad y_n] = \begin{bmatrix} x_1y_1 & x_1y_2 & \dots & x_1y_n \\ x_2y_1 & x_2y_2 & \dots & x_2y_n \\ \vdots & \vdots & & \vdots \\ x_ny_1 & x_ny_2 & \dots & x_ny_n \end{bmatrix} \in \mathbb{R}^{m \times n}$$

Vector Norm

- Length of vector
- Examples
 - 1-norm:
 $\|x\|_1 = \sum_{i=1}^n |x_i|$
 - 2-norm:
 $\|x\|_2 = \sqrt{\sum_{i=1}^n |x_i|^2}$
 - ∞ -norm:
 $\|x\|_\infty = \max(|x_1|, |x_2|, \dots, |x_n|)$

p-norm

$$\|x\|_p = \left(\sum_{i=1}^n |x_i|^p \right)^{\frac{1}{p}}$$

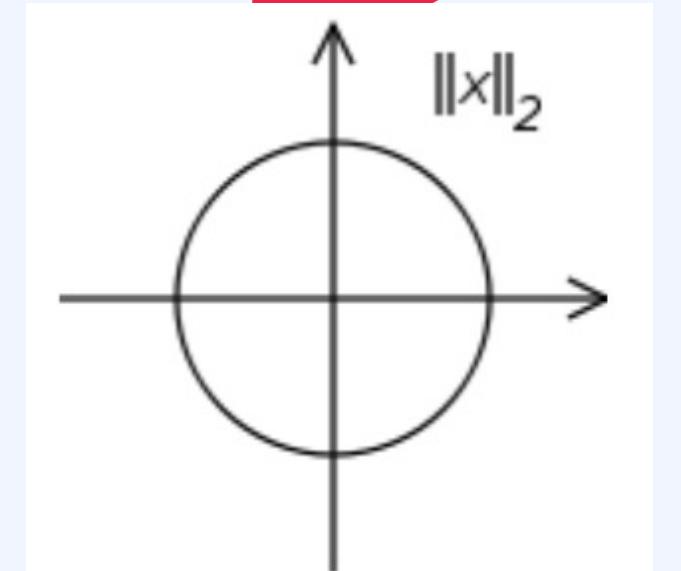


Vector Norm

- 2-norm revisited

$$\mathbf{x} = [x_1 \quad x_2 \quad \dots \quad x_n]^\top \in \mathbb{R}^n$$

$$\mathbf{y} = [y_1 \quad y_2 \quad \dots \quad y_n]^\top \in \mathbb{R}^n$$

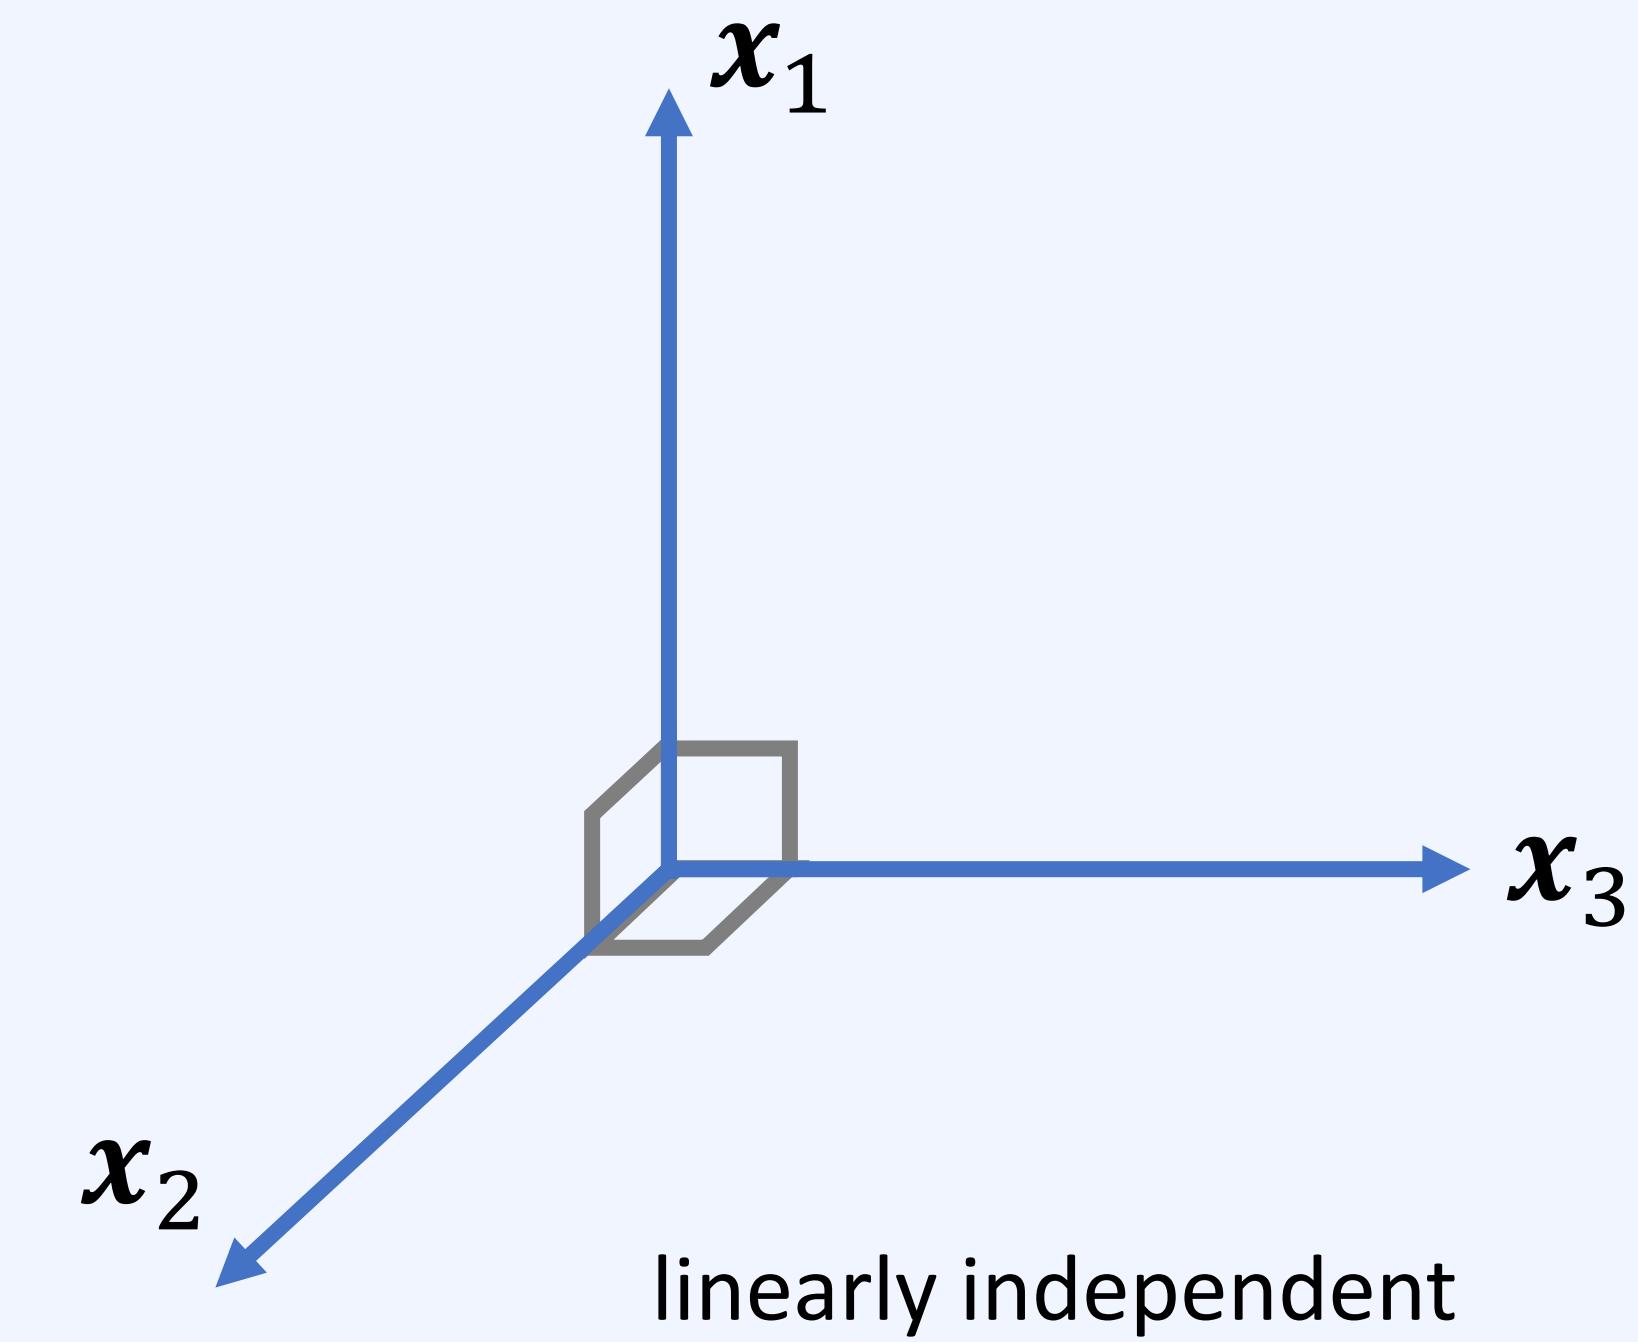
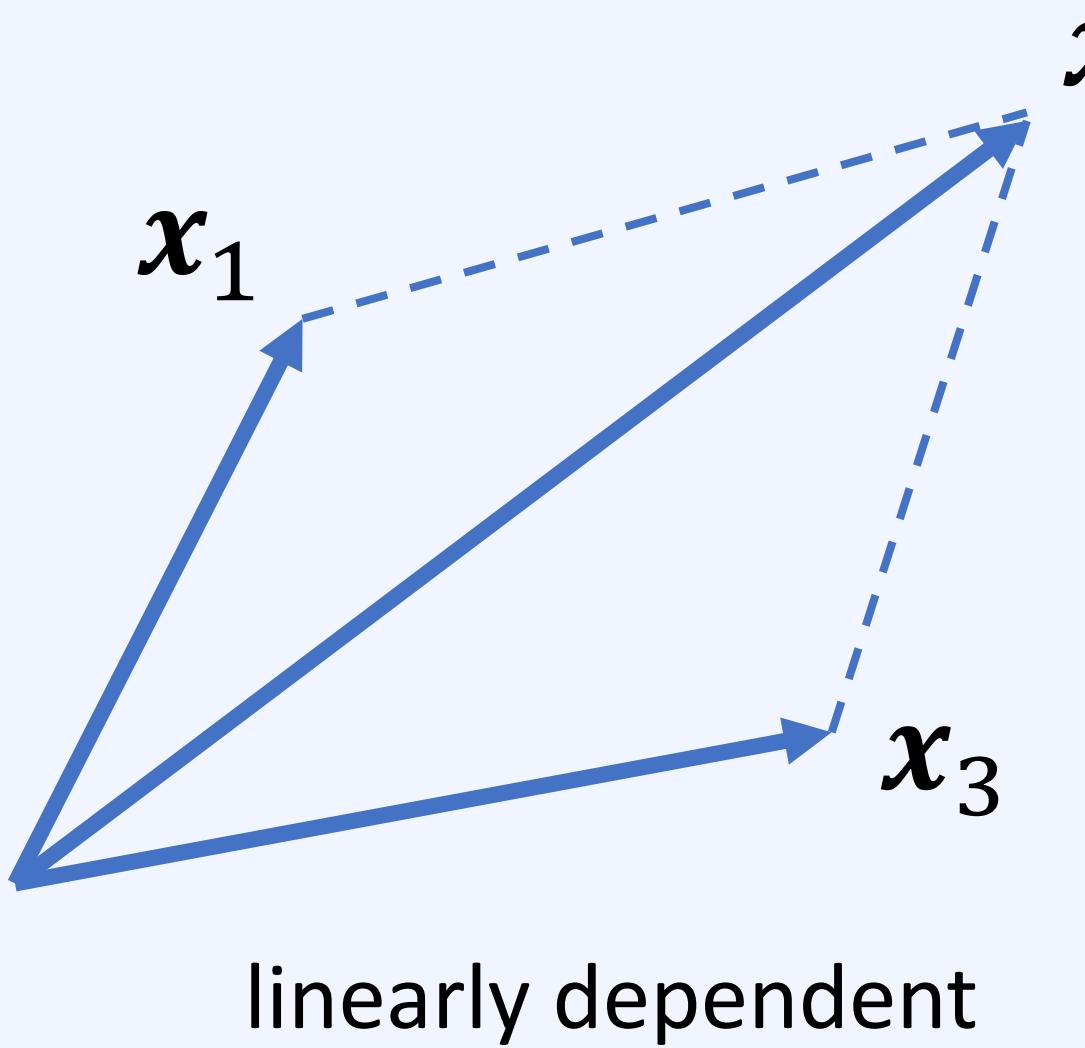


$$\|\mathbf{x}\|_2^2 = \sum_{i=1}^n |x_i|^2 = x_1^2 + x_2^2 + \dots + x_n^2 = \mathbf{x}^\top \mathbf{x} = \mathbf{x} \cdot \mathbf{x}$$

$$\|\mathbf{x} - \mathbf{y}\|_2^2 = \sum_{i=1}^n |x_i - y_i|^2 = (\mathbf{x} - \mathbf{y})^\top (\mathbf{x} - \mathbf{y}) = \|\mathbf{x}\|_2^2 - 2\mathbf{x}^\top \mathbf{y} + \|\mathbf{y}\|_2^2$$

Linear Dependency

- Given a set of vectors $X = \{x_1, x_2, \dots, x_k\}$
- Linear combination** of vectors: $a_1x_1 + a_2x_2 + \dots + a_kx_k$
- $x_i \in X$ is **linearly dependent** if it can be represented by a linear combination of $X \setminus \{x_i\}$.



Basis

- A linearly independent set of vectors that spans the whole space.
- Standard basis (a.k.a. unit vectors)

$$\mathbf{e}_1 = \begin{bmatrix} 1 \\ 0 \\ \vdots \\ 0 \end{bmatrix}, \mathbf{e}_2 = \begin{bmatrix} 0 \\ 1 \\ \vdots \\ 0 \end{bmatrix}, \dots, \mathbf{e}_n = \begin{bmatrix} 0 \\ 0 \\ \vdots \\ 1 \end{bmatrix}$$

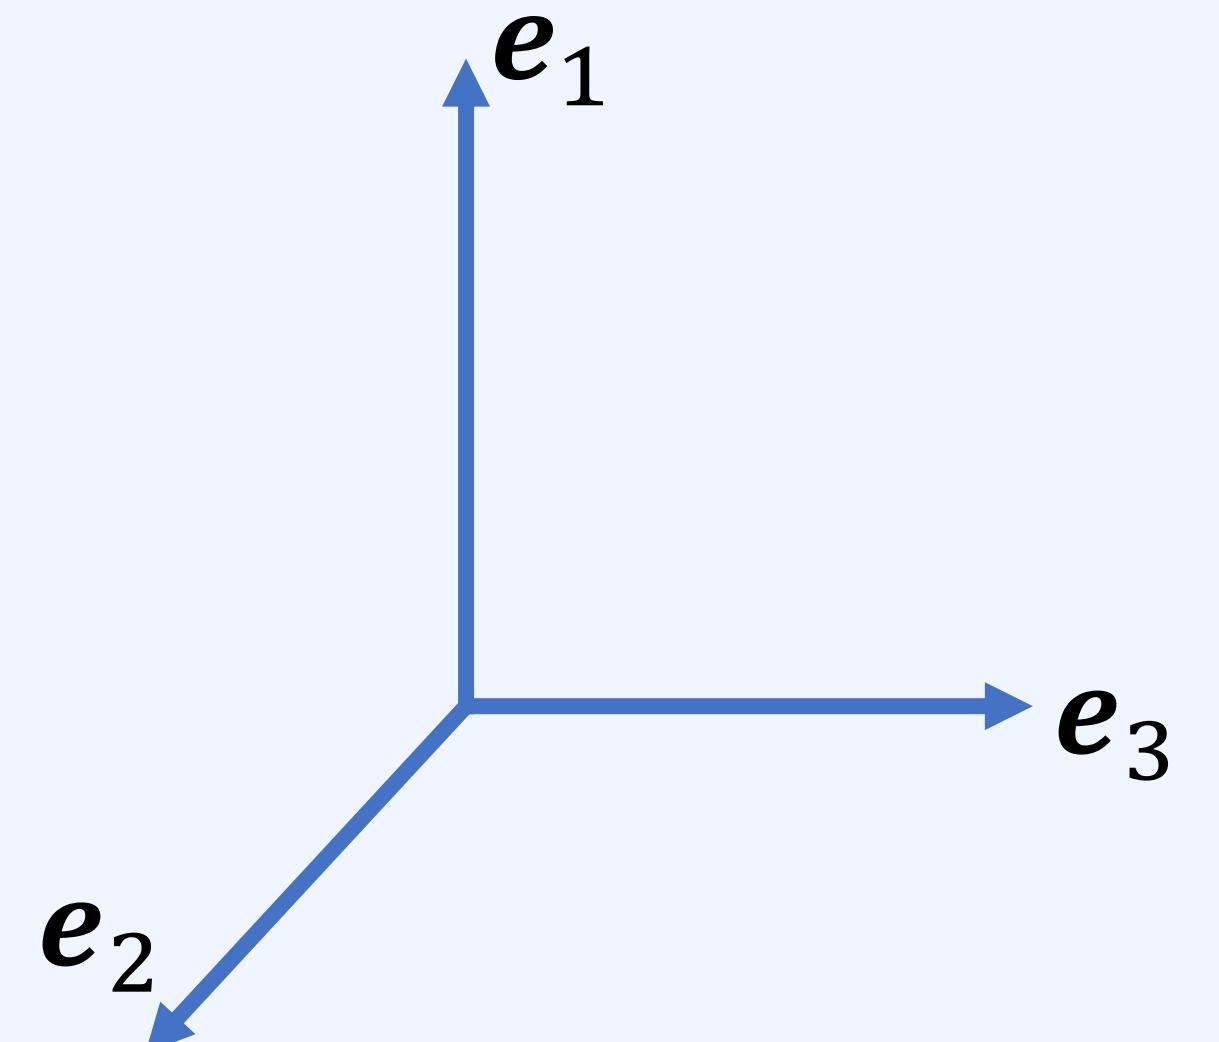
- Example: $[5 \ 2 \ 3]^\top = 5\mathbf{e}_1 + 2\mathbf{e}_2 + 3\mathbf{e}_3$
- Projection of a vector:

$$\mathbf{x} \cdot \mathbf{e}_i = \mathbf{x}^\top \mathbf{e}_i$$

$$\mathbf{x} = [\mathbf{x}^\top \mathbf{e}_1 \ \mathbf{x}^\top \mathbf{e}_2 \ \dots \ \mathbf{x}^\top \mathbf{e}_n]$$

- Orthogonal: $\mathbf{e}_i^\top \mathbf{e}_j = 0, \quad \forall i \neq j$
- Normalized: $\mathbf{e}_i^\top \mathbf{e}_i = 1$

} **Orthonormal**



Matrix

- Rectangular array of numbers
- Notation (표기법)

$$X \in \mathbb{R}^{m \times n}. \quad X = \begin{bmatrix} x_1y_1 & x_1y_2 & \dots & x_1y_n \\ x_2y_1 & x_2y_2 & \dots & x_2y_n \\ \vdots & \vdots & & \vdots \\ x_ny_1 & x_ny_2 & \dots & x_ny_n \end{bmatrix}. \quad [X]_{ij} = x_{ij}.$$

- Special matrices
 - Square matrix: matrix with (# of rows) == (# of columns)
 - Identity matrix: matrix with diagonal elements ones and non-diagonal elements zeros,

$$XI = IX = X$$

Matrix Operations (행렬 연산과, 그 연산들의 성질)

- Addition (덧셈)
 - Commutative(교환법칙): $X + Y = Y + X$
 - Associative(결합법칙): $(X + Y) + Z = X + (Y + Z)$
- Subtraction (뺄셈)
 - Not-commutative, not-associative
- Multiplication (곱셈)
 - Associative: $(XY)Z = X(YZ)$
 - Distributive: $X(Y + Z) = XY + XZ$
 - Non-commutative(결합법칙 성립안함): $XY \neq YX$
- Transpose: $[X^T]_{ij} = [X]_{ji}$
 - $(X^T)^T = X$
 - $(XY)^T = Y^T X^T$

$$\begin{bmatrix} 3 & 8 \\ 4 & 6 \end{bmatrix} + \begin{bmatrix} 4 & 0 \\ 1 & -9 \end{bmatrix} = \begin{bmatrix} 7 & 8 \\ 5 & -3 \end{bmatrix}$$

$3+4=7$

$$\begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix} \times \begin{bmatrix} 7 & 8 \\ 9 & 10 \\ 11 & 12 \end{bmatrix} = \begin{bmatrix} 58 \end{bmatrix}$$

"Dot Product"

$$\begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix} \times \begin{bmatrix} 7 & 8 \\ 9 & 10 \\ 11 & 12 \end{bmatrix} = \begin{bmatrix} 58 & 64 \end{bmatrix}$$

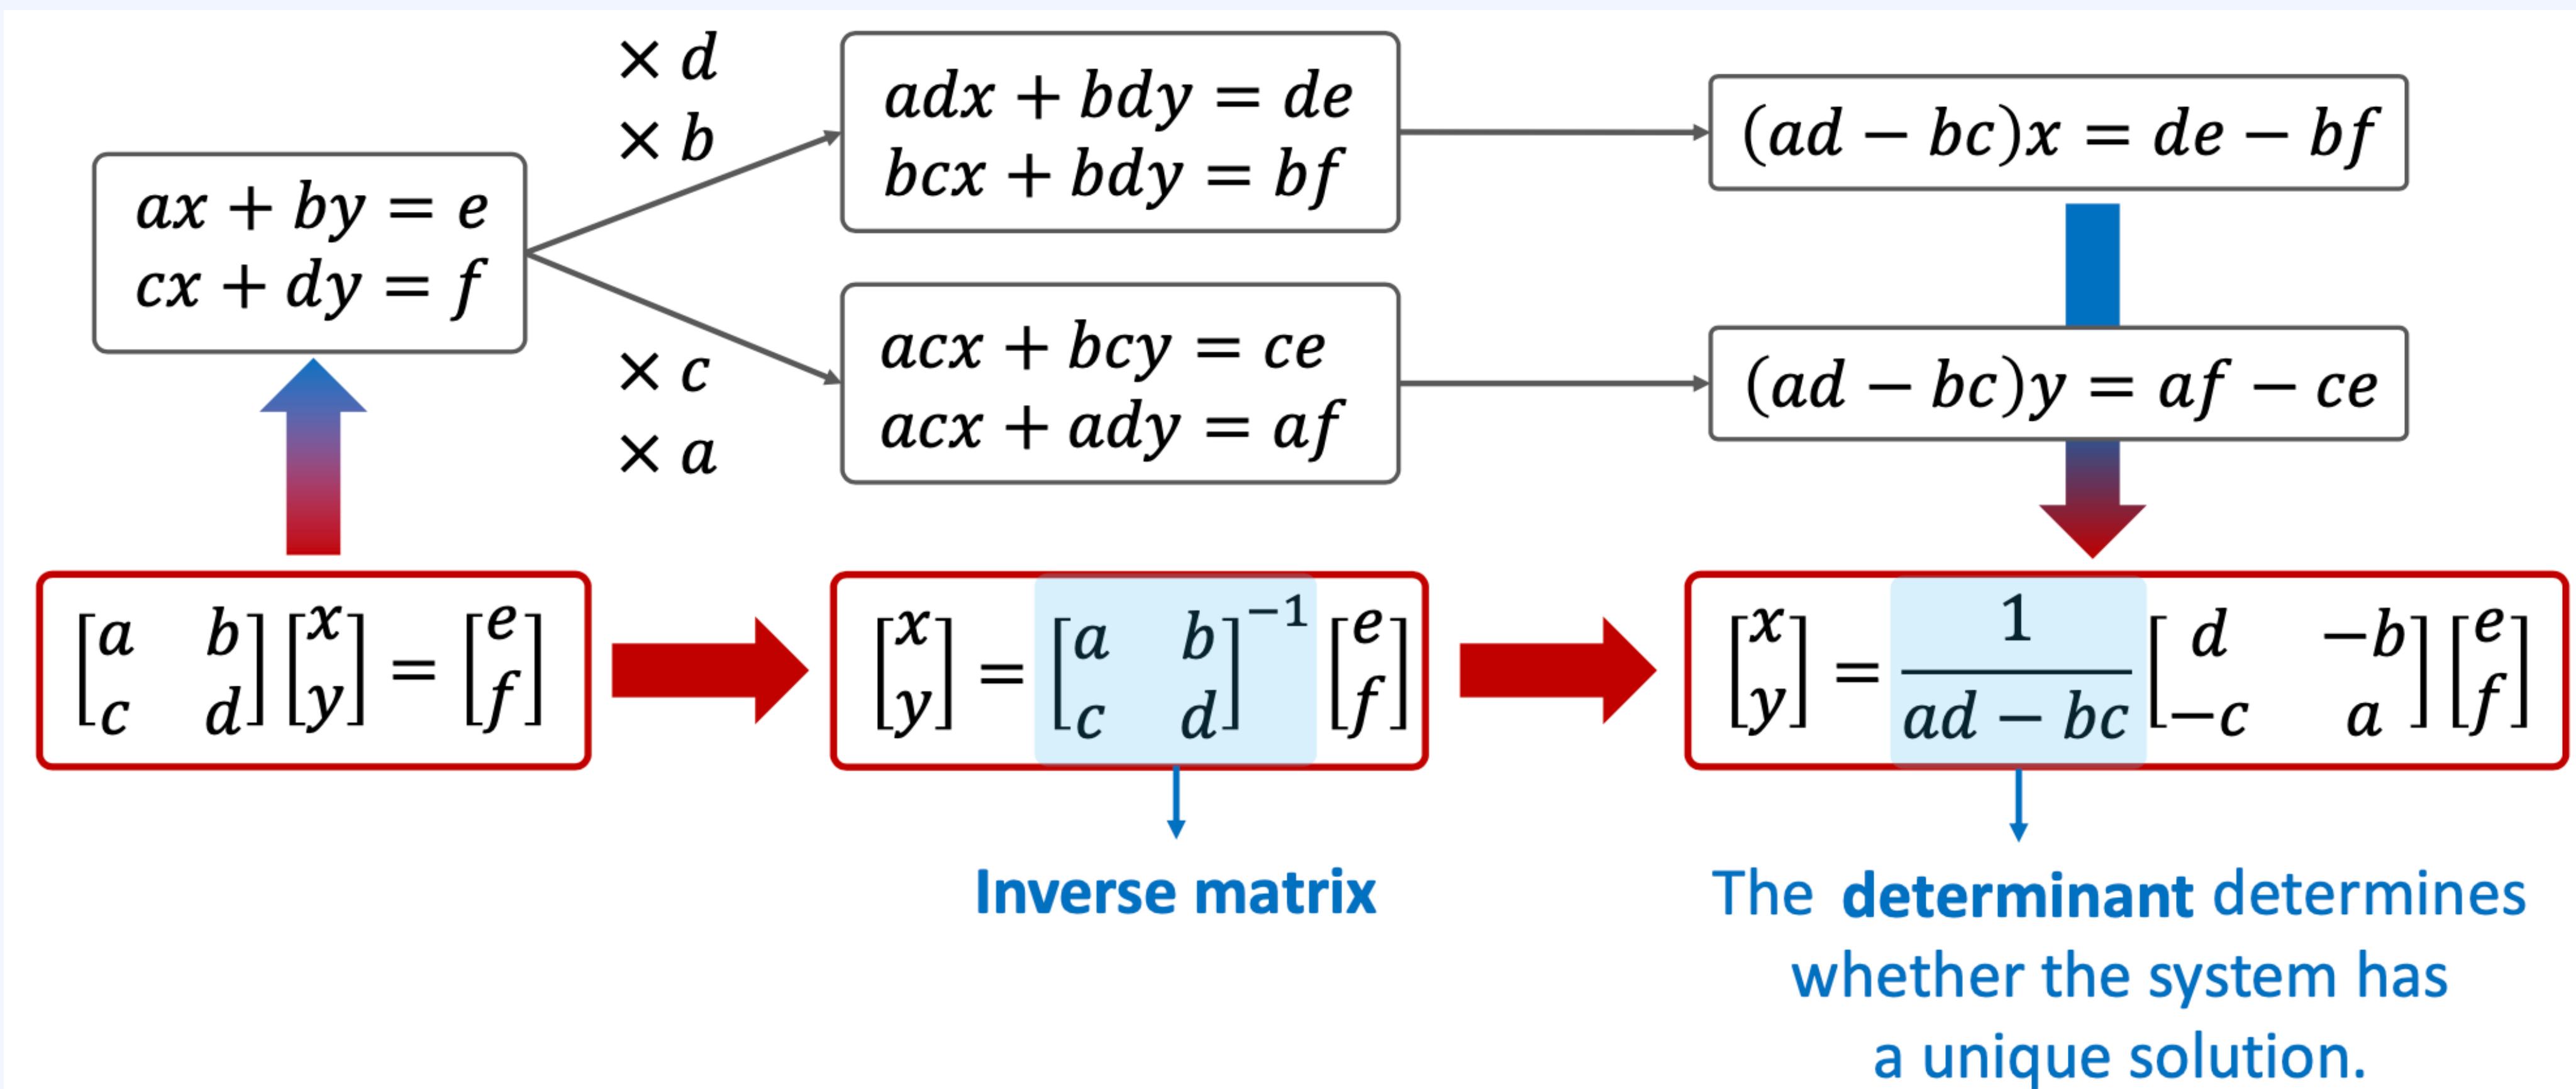
$$\begin{bmatrix} 6 & 4 & 24 \\ 1 & -9 & 8 \end{bmatrix}^T = \begin{bmatrix} 6 & 1 \\ 4 & -9 \\ 24 & 8 \end{bmatrix}$$

Rank of Matrix

- Definition
 - The number of **linearly independent** rows or columns in a matrix
- Examples
 - $\text{rank}\begin{pmatrix} -1 & 2 \\ 4 & -8 \end{pmatrix} = 1.$ $\text{rank}\begin{pmatrix} -5 & -3 \\ 7 & 2 \end{pmatrix} = 2$
- Properties
 - $X \in \mathbb{R}^{m \times n}$, $\text{rank}(X) \leq \min(m, n)$
 - X is full-rank if $\text{rank}(X) = \min(m, n)$
 - $\text{rank}(X) = \text{rank}(X^\top)$
 - $\text{rank}(XY) \leq \min(\text{rank}(X), \text{rank}(Y))$
 - $\text{rank}(X + Y) \leq \text{rank}(X) + \text{rank}(Y)$
 - A square matrix with full rank is **non-singular**;
 - otherwise, the matrix is **singular**.

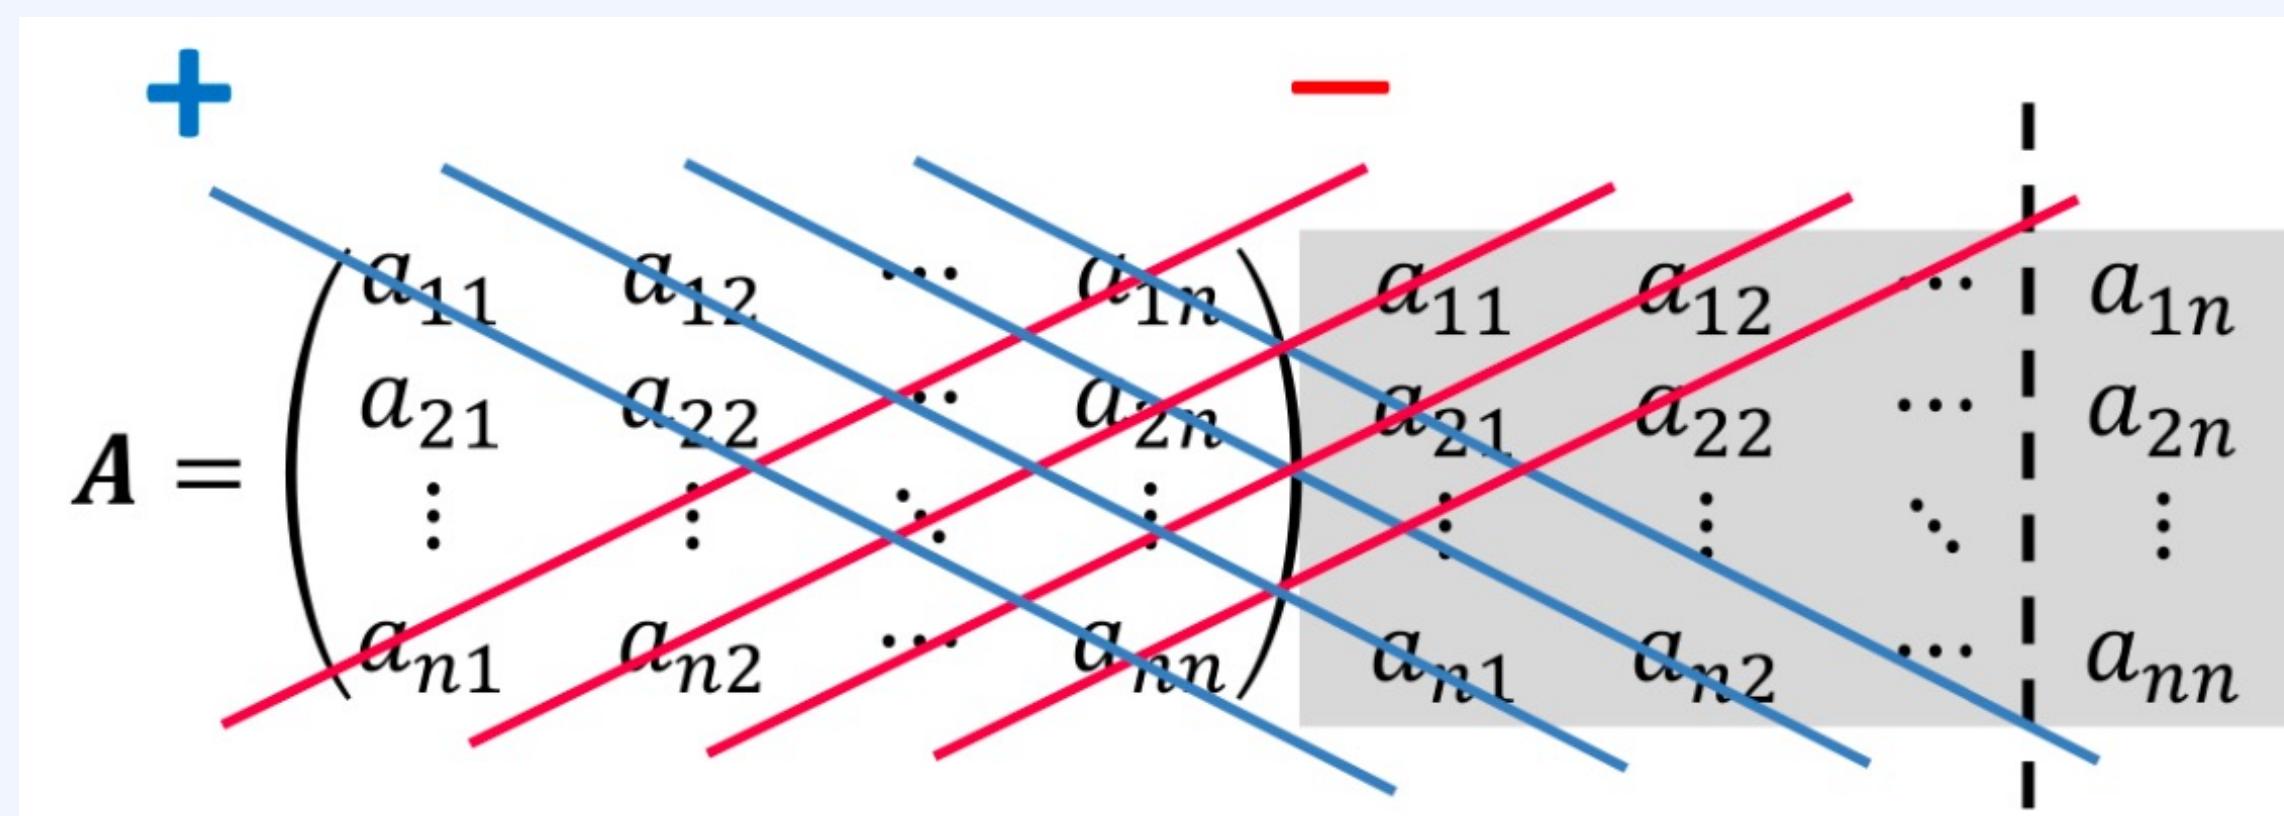
Square Matrix and Linear Equations

- Solving linear equations given by a square matrix



Determinant

- How to compute



MATLAB command: `det(A)`
Python command: `numpy.linalg.det(A)`

- Properties
 - Computed from the elements of a square matrix
 - $\det(A) = 0$ iff A is singular.
 - $\det(AB) = \det(A)\det(B)$
 - $\det(A^{-1}) = \det(A)^{-1}$
 - $\det(\lambda A) = \lambda^n \det(A)$

Inverse Matrix

- Condition to have inverse matrix: **square** and **non-singular**
- Inverse matrix of X : X^{-1}
 - $XX^{-1} = X^{-1}X = I$
- 2D example
 - $\begin{bmatrix} a & b \\ c & d \end{bmatrix} = \frac{1}{ad-bc} \begin{bmatrix} d & -b \\ -c & a \end{bmatrix}$
- Properties
 - $(XY)^{-1} = Y^{-1} X^{-1}$
 - $X^{-1} = X^{-1}$
 - $(X^\top)^{-1} = (X^{-1})^\top$
- For an **orthonormal matrix**, $X^{-1} = X^\top$
- For a diagonal matrix $D = \text{diag}(d_1, d_2, \dots, d_n) \in \mathbb{R}^{m \times n}$
 - $D^{-1} = \text{diag}(d_1^{-1}, d_2^{-1}, \dots, d_n^{-1})$

Orthonormal matrix

A square matrix with real entries whose columns and rows are orthonormal vectors.

- Identity matrix: $\begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$
- Rotation matrix: $\begin{bmatrix} \cos(\theta) & -\sin(\theta) \\ \sin(\theta) & \cos(\theta) \end{bmatrix}$

Determinant, Inverse Matrix, and Linear Equations

- $\det \begin{pmatrix} a & b \\ c & d \end{pmatrix} = ad - bc = 0$

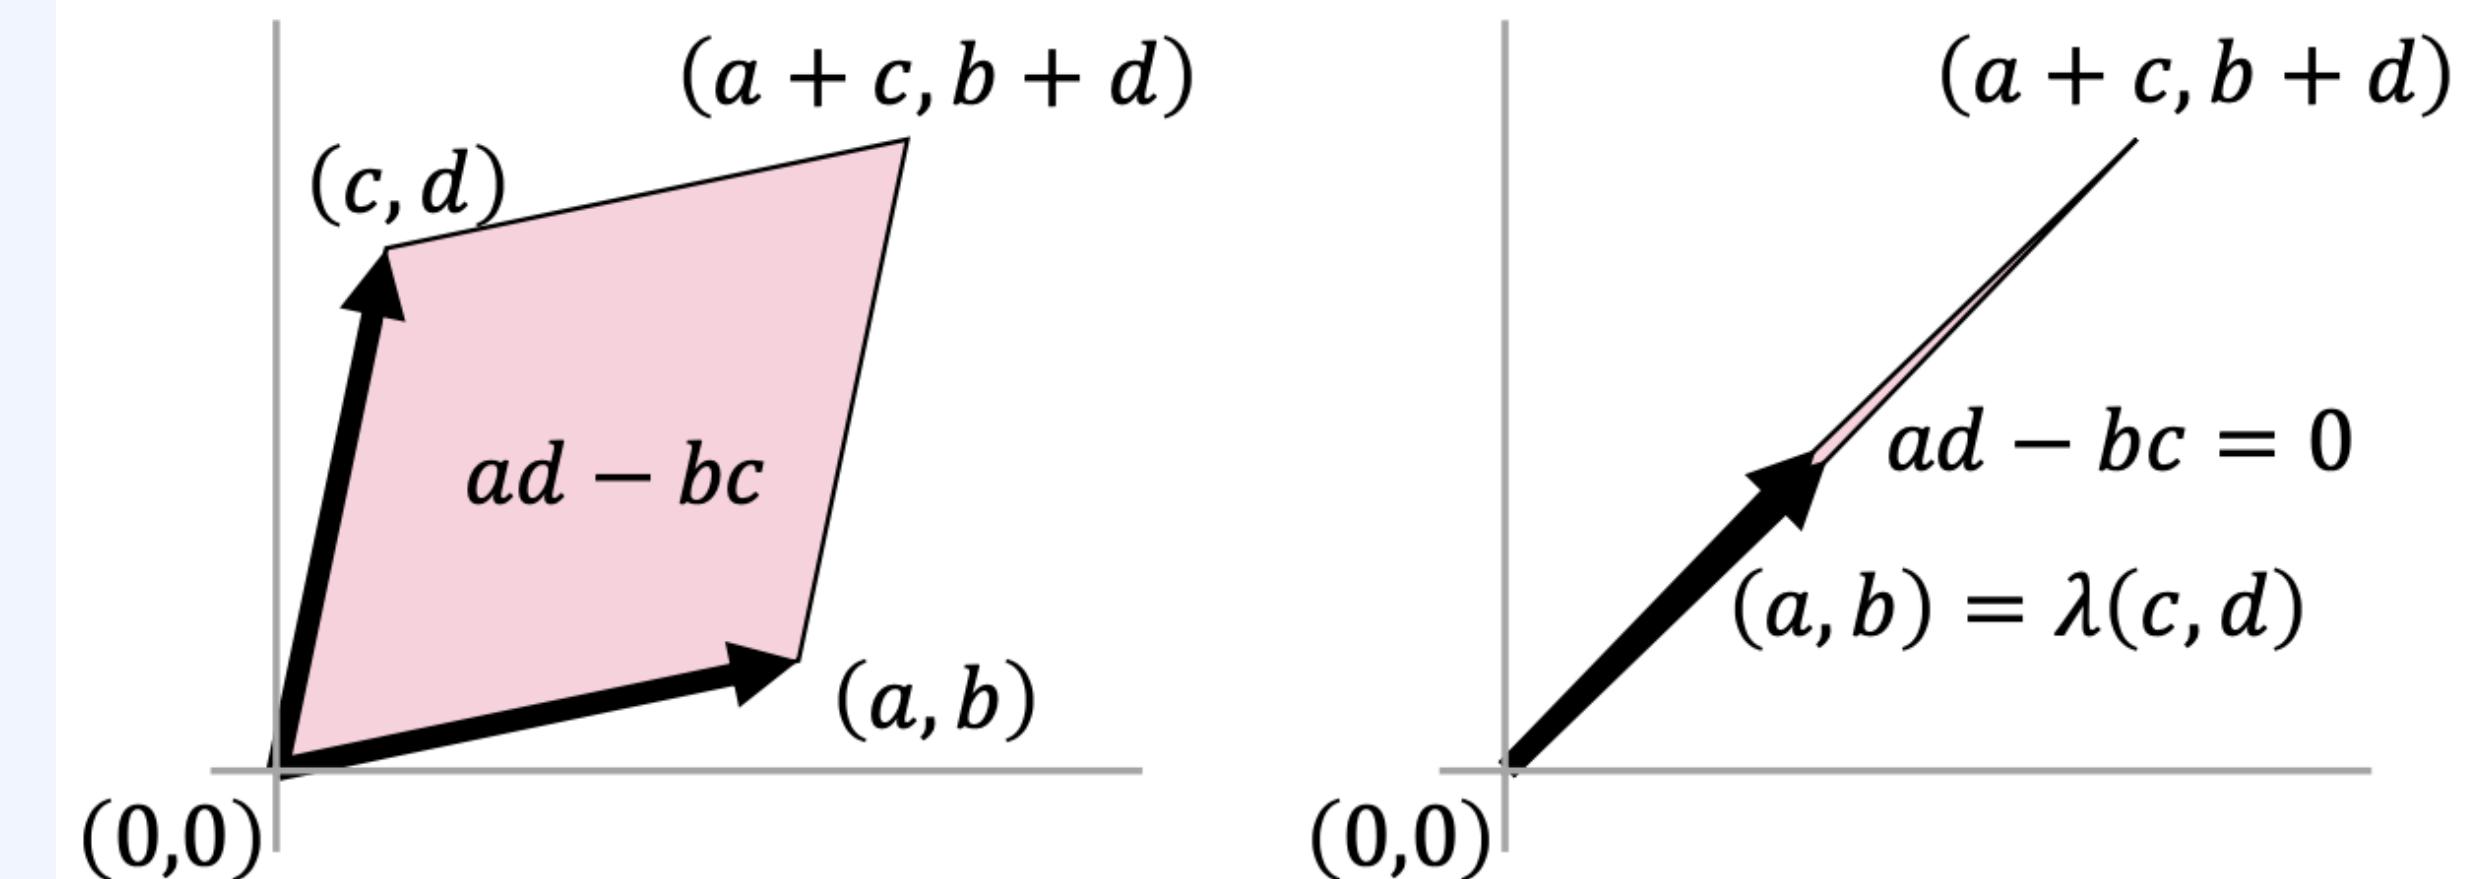
- Geometric interpretation
(singular matrix)

- No inverse matrix

$$\begin{bmatrix} a & b \\ c & d \end{bmatrix}^{-1} = \frac{1}{ad - bc} \begin{bmatrix} d & -b \\ -c & a \end{bmatrix}$$

- No solution to the linear equations

$$0 \cdot x = de - bf, \quad 0 \cdot y = af - ce$$



Eigen Decomposition

- Eigenvector and eigenvalue
 - Defined for a square matrix A
 - Non-zero vector x for which there is a $\lambda \in \mathbb{R}$ such that $Ax = \lambda x$
 - Eigenvector: x
 - Eigenvalue: λ
- How to compute
 - Solve $(A - \lambda I)x = 0$ for x and λ
- Simply,
 - MATLAB command: $[V, D] = \text{eig}(A)$
 - Python command: $[V, D] = \text{np.linalg.eig}(A)$

Eigen Decomposition

- A square and symmetric matrix A can be decomposed as

$$A = VDV^T = [v_1 \ v_2 \ \dots \ v_n] \begin{bmatrix} \lambda_1 & 0 & \dots & 0 \\ 0 & \lambda_2 & \dots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & \lambda_n \end{bmatrix} [v_1 \ v_2 \ \dots \ v_n]$$

$$\begin{aligned} v_i^T v_i &= 1 \\ v_i^T v_j &= 0 \end{aligned}$$

- V is an **orthonormal matrix** of A 's eigenvectors.
- D is an **diagonal matrix** of the associated eigenvalues.
- Properties
 - A square matrix is singular if any of its eigenvalues is zero.
 - A square matrix A is positive definite (negative definite) if all eigenvalues are positives (negatives).
 - Semi-definiteness: $\forall x \neq 0, \quad x^T A x > 0 \quad (x^T A x < 0)$
 - $\forall x \neq 0, \quad x^T A x \geq 0 \quad (x^T A x \leq 0)$

Eigen Decomposition

- Computing inverse matrix
 - Problem: $AX = I$
 - Decomposition: $A = VDV^T$
 - Problem revisited: $A = VDV^TX = I$
 - Solution: $X = VD^{-1}V^T$
- Solving optimization problems
 - Objective function: $\min\|Ax\|^2$
 - Find the eigenvector of A^TA corresponding to the least eigenvalue (e.g., PCA)

$$\min\|Ax\|^2 = \min\|\lambda x\|^2$$

The least eigenvalue

Corresponding eigenvector

If you need more about linear algebra,

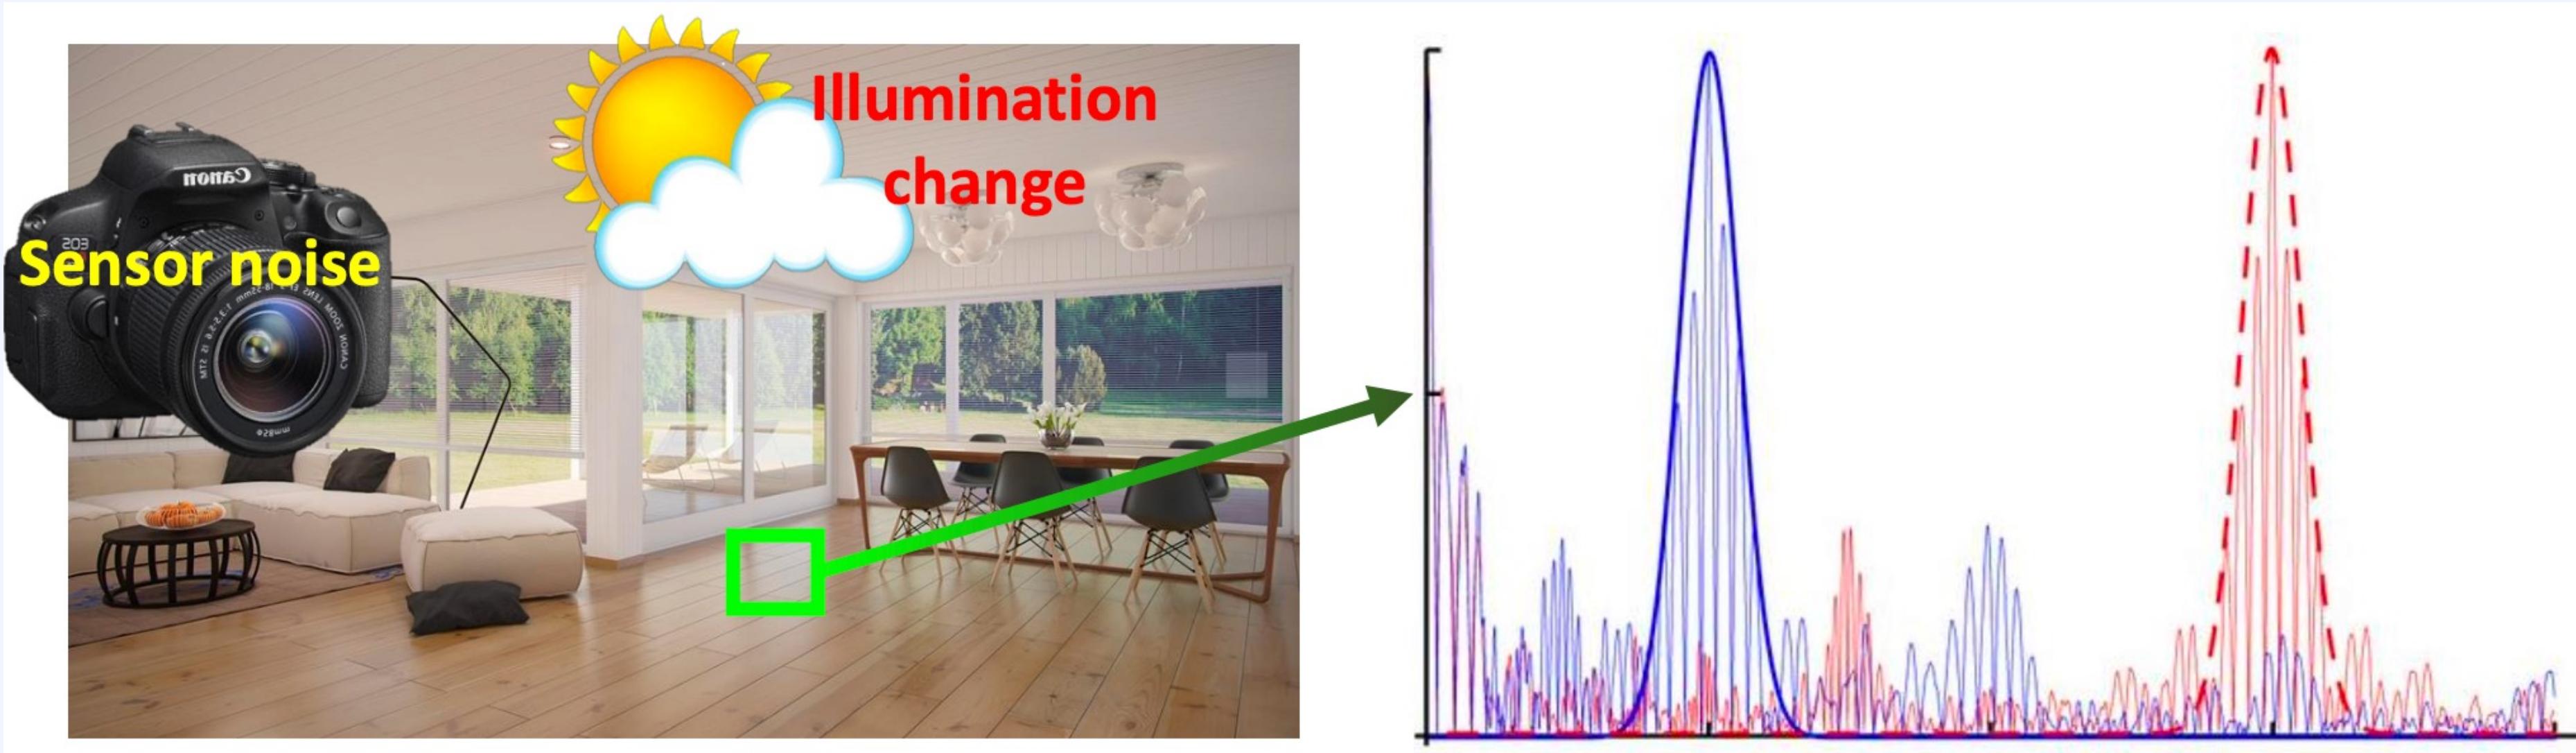
- Matrix Cookbook
 - <https://www.math.uwaterloo.ca/~hwolkowi/matrixcookbook.pdf>

컴퓨터 비전 초격차: 이론/실습

3 Preliminary – Probability

Probability in Computer Vision

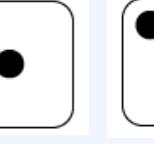
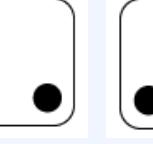
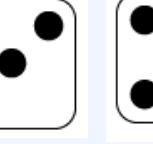
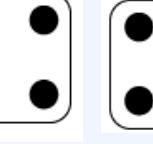
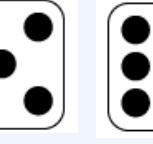
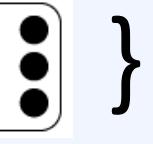
- Visual sensing is noisy and uncertain.



- Probabilistic models are useful to handle such uncertainty.

Most of the slides are inspired by POSTECH
AIGS549 class (Prof. Suha Kwak)

Random variable

- Sample Space (Ω)
 - The set of all possible outcomes to an experiment
 - Example: Tossing a single six-sided dice: { , , , , ,  }
- Random Variable ($X: \Omega \rightarrow \mathbb{R}$)
 - A function that assigns a (usually *real*) number to each point in a sample space Ω
 - A variable whose possible values are numerical outcomes of a random phenomenon
 - Discrete or continuous
 - Examples:

$$X(\text{••}) = 3$$

$$X(\text{•••}) = 5$$

Discrete
Random Variable

The probability that X takes value ≤ 3

$$= P(\omega \in \Omega : X(\omega) \leq 3)$$

$$= P(X \leq 3)$$

Probability Axioms

- Sample space (Ω): The set of all possible outcomes
- Event space (E): A power set whose element is a set of Ω
 - If $A \in E$, then $A \subseteq \Omega$
- Probability measure (P)
 - A function $P: E \rightarrow \mathbb{R}$ with the following conditions:
 - $P(A) \geq 0, \forall A \in E$
 - $P(\Omega) = 1$
 - $\sum_i P(A_i) = P(A_1 \cup A_2 \cup A_3 \cup \dots)$, if A_i 's are disjoint

Important Properties in Probability

- Set inclusion
 - $A \subseteq B \Rightarrow P(A) \leq P(B)$
- Union
 - $P(A \cup B) = P(A) + P(B) - P(A \cap B)$
- Complement
 - $P(\Omega \setminus A) = 1 - P(A)$
- Independence
 - If A and B are independent, $P(A \cap B) = P(A)P(B)$
- Notation
 - $P(A \cap B) = P(A, B)$

Conditional Probability

- Conditional probability of A given B

- $P(A|B) = \frac{P(A,B)}{P(B)}$

- $P(A,B) = P(A|B) P(B)$

- Independence

$$P(A|B) = \frac{P(A,B)}{P(B)} = \frac{P(A)P(B)}{P(B)} = P(A)$$

- Conditional independence

- A and B are conditionally independent given C.

$$P(A, B|C) = P(A|C)P(B|C)$$

$$P(A|B, C) = P(A|C)$$

$$P(\text{rain, sunny})$$

$$= P(\text{sunny}) P(\text{rain|sunny})$$

$$= 0.5 \times 0.2 = 0.1$$

$$P(\text{rain, cloudy})$$

$$= P(\text{cloudy}) P(\text{rain|cloudy})$$

$$= 0.5 \times 0.7 = 0.35$$

$$P(T, B)$$

$$= P(T) P(B|T)$$

$$= P(T) P(B)$$

$$= 0.5 \times 0.5$$

Bayes' Theorem

- Revisit conditional probabilities

$$P(A|B) = \frac{P(A,B)}{P(B)}, P(B|A) = \frac{P(A,B)}{P(A)}$$

$$P(A, B) = P(A|B) P(B) = P(B|A) P(A)$$

- Bayes' Theorem

$$P(A|B) = \frac{P(B|A)P(A)}{P(B)}$$

posterior \propto likelihood \times prior

Bayes' Theorem

- Example

posterior \propto *likelihood* \times *prior*

$$P(A|B) = \frac{P(B|A)P(A)}{P(B)}$$

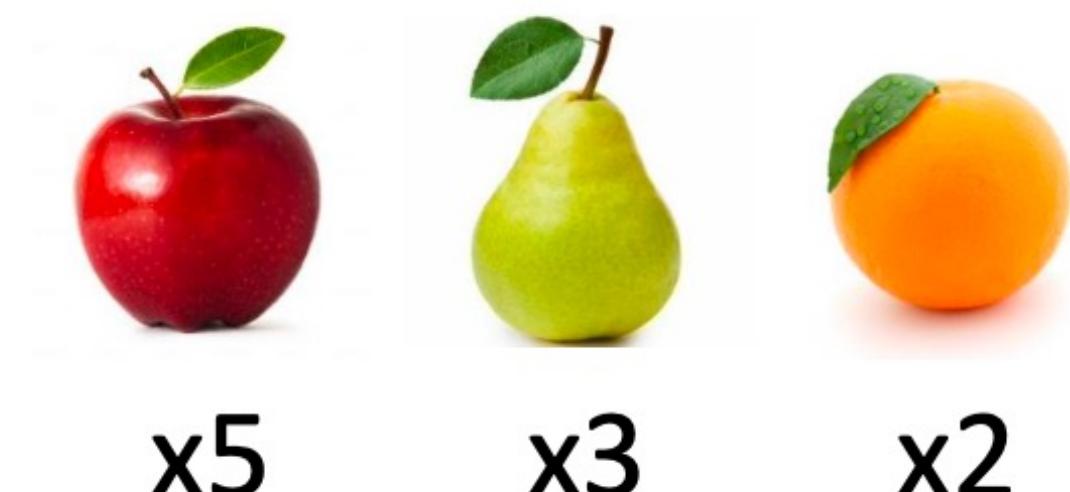
[*prior*]

$$P(\text{apple}) = 0.5, \quad P(\text{pear}) = 0.3, \quad P(\text{orange}) = 0.2$$



[*likelihood*]

$$p(\text{tongue} | \text{apple}) = 0.3, \quad p(\text{tongue} | \text{pear}) = 0.5, \quad p(\text{tongue} | \text{orange}) = 0.9$$



[*posterior* \propto *likelihood* \times *prior*]

$$P(\text{apple} | \text{tongue}) \propto 0.15, \quad P(\text{pear} | \text{tongue}) \propto 0.15, \quad P(\text{orange} | \text{tongue}) \propto 0.18$$

Gaussian Distributions

- Univariate Gaussian distribution

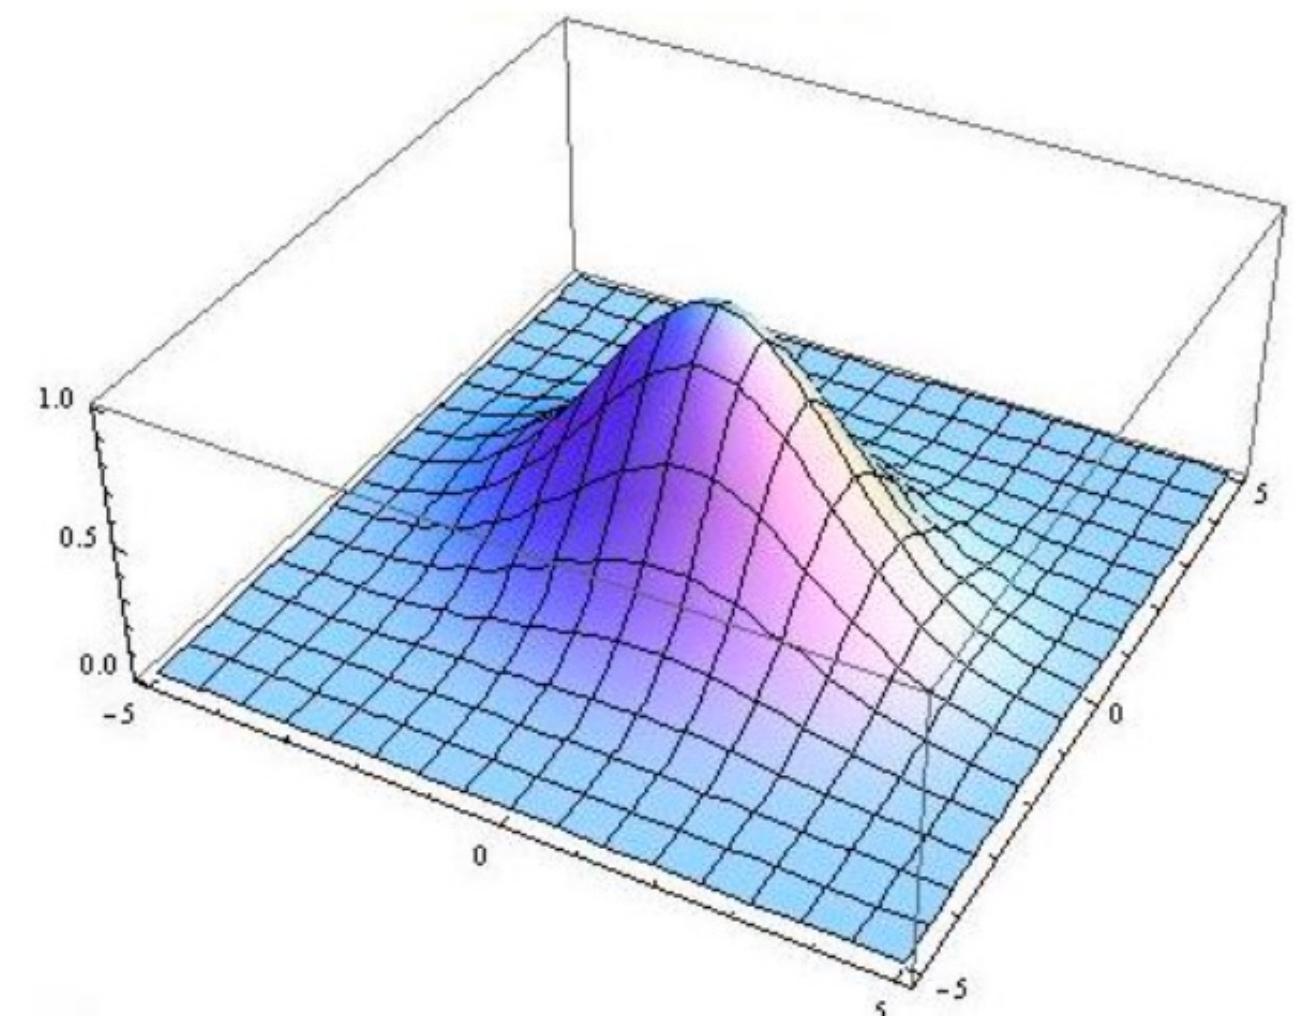
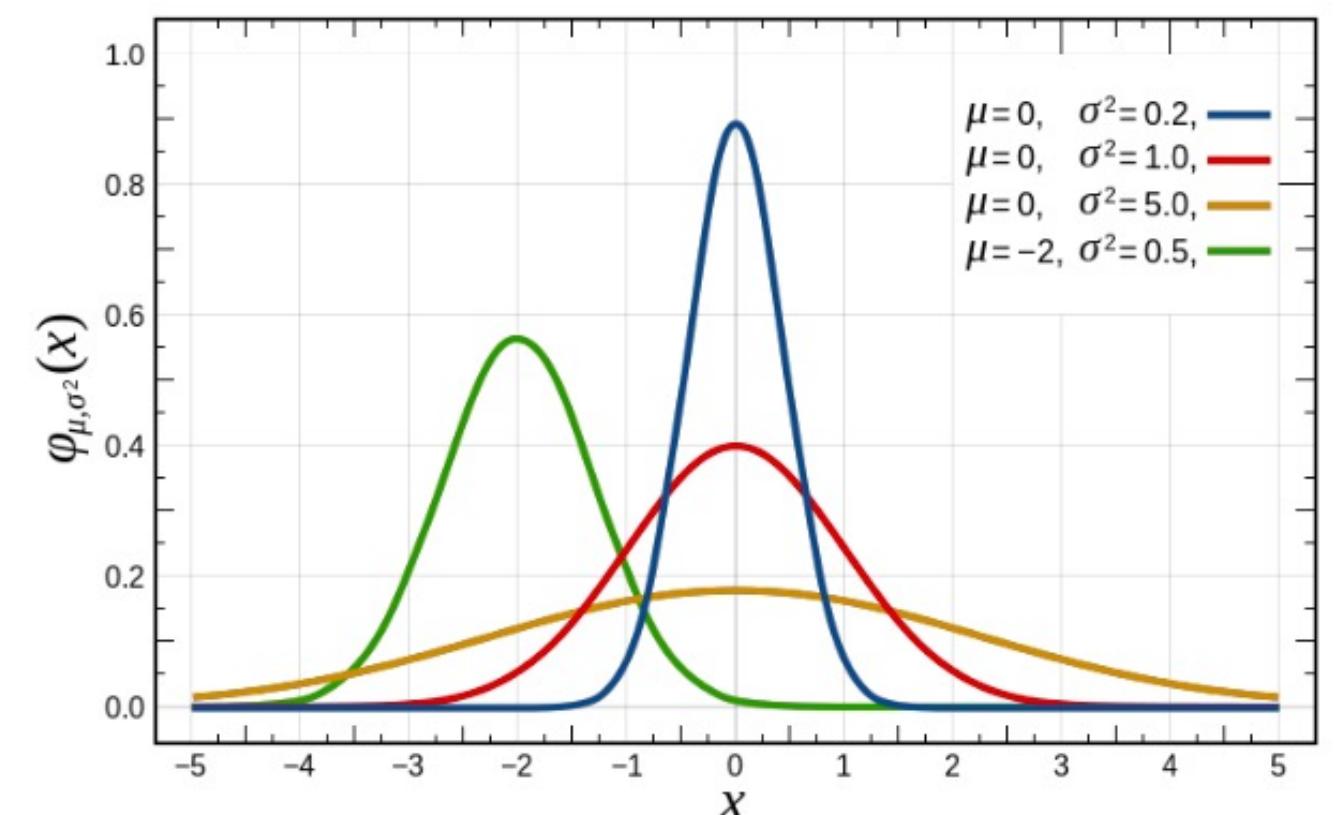
$$f(x; \mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x - \mu)^2}{2\sigma^2}\right)$$

- Multivariate Gaussian distribution

$$f(\mathbf{x}; \boldsymbol{\mu}, \boldsymbol{\Sigma}) = \frac{1}{(2\pi)^{d/2} |\boldsymbol{\Sigma}|^{1/2}} \exp\left(-\frac{1}{2} (\mathbf{x} - \boldsymbol{\mu})^\top \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu})\right),$$

where $\mathbf{x} = [x_1, x_2, \dots, x_d]^\top$, $\boldsymbol{\mu} = [\mu_1, \mu_2, \dots, \mu_d]^\top$ and

$$\boldsymbol{\Sigma} = \begin{bmatrix} \text{cov}(x_1, x_1) & \text{cov}(x_1, x_2) & \cdots & \text{cov}(x_1, x_d) \\ \text{cov}(x_2, x_1) & \text{cov}(x_2, x_2) & \cdots & \text{cov}(x_2, x_d) \\ \vdots & \vdots & \ddots & \vdots \\ \text{cov}(x_d, x_1) & \text{cov}(x_d, x_2) & \cdots & \text{cov}(x_d, x_d) \end{bmatrix}.$$



Why Gaussian?

- Only two parameters to estimate
 - Mean μ and covariance matrix Σ
- Self-conjugate
 - Gaussian \times Gaussian = Gaussian
 - Bayes' theorem

컴퓨터 비전 초격차: 이론/실습

4 Preliminary – OpenCV, PIL, matplotlib

Index

- OpenCV
- matplotlib
- PIL
- Scikit-image
- Example codes

Most of the slides are inspired by POSTECH
AIGS549 class (Prof. Suha Kwak)

OpenCV

- 컴퓨터 비전을 목적으로 하는 오픈 소스 라이브러리
- 인텔 CPU에서 사용하는 경우 속도 향상을 볼 수 있는 IPP (Intel Performance Primitives)를 지원한다.
- 기존에 C++에서 사용할 수 있게 구현되었으나, OpenCV-python을 통해 python 포팅도 되어 있다.
- 단점: **GPU operation**에 대한 지원을, 명시적으로 python과 연계하여 하지는 않는다.
- 설치: pip install opencv-python
- OpenCV documentation: <https://docs.opencv.org/4.x/>
- OpenCV python tutorial:
https://docs.opencv.org/4.x/d6/d00/tutorial_py_root.html

Matplotlib

- Python 과 numpy array를 기반으로 plotting 과 visualize를 목적으로 하는 라이브러리
- 주어진 데이터에 대해서 차트와 plot 을 편리하게 그려주는 데이터 시각화 패키지 이다.
- 우리는 본 강의에서 이미지 시각화와 디버깅을 위해 사용한다.
- 설치 : pip install matplotlib
- Matplotlib documentation : <https://matplotlib.org/stable/index.html>
- Matplotlib tutorials: <https://matplotlib.org/stable/tutorials/index.html>

OpenCV

- Example 1 : Image read and visualize
- Example 2 : multiple image visualize, figure/axs in matplotlib
- Example 3 : Image cropping, masking and save
- Example 4 : Image matching and visualize

PIL

Preliminary

Linear Algebra
Basic

- PIL (Python Imaging Library)로서 컴퓨터 비전 보다는 이미지 처리에 중점을 둔 라이브러리.
- 픽셀단위의 이미지 조작이나, 마스킹, 투명도 제어, 윤곽 보정 및 검출 등의 다양한 이미지 조작을 할 수 있다.
- 특히,
- 설치 : pip install pillow
- Pillow (PIL forked) documentation:
<https://pillow.readthedocs.io/en/stable/>

PIL

Preliminary

Linear Algebra
Basic

- Image read and visualize
- Image cropping, rotating, scaling
- Image interpolation (upsampling, downsampling)

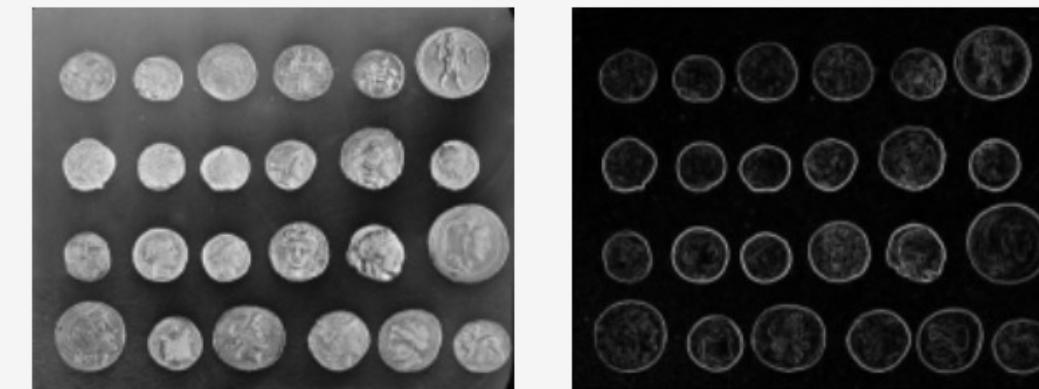
Scikit-image

- Scikit-image 는 Pillow (PIL) 과 마찬가지로, 이미지 조작과 필터링이 가능하다.
- numpy 를 기반으로 동작하기 때문에, 좀 더 numpy와의 호환성이 좋다.

Getting Started

Filtering an image with `scikit-image` is easy! For more examples, please visit our [gallery](#).

```
from skimage import data, io, filters
image = data.coins()
# ... or any other NumPy array!
edges = filters.sobel(image)
io.imshow(edges)
io.show()
```



- 설치 : pip install scikit-image
- Documentation: <https://scikit-image.org/>

Environment

Preliminary

Linear Algebra
Basic

- Jupyter notebook
- If you don't have GPU resources, please use Google CoLab!!

컴퓨터 비전 초격차: 이론/실습

5 Preliminary – PyTorch, Torchvision, conda

Index

- PyTorch
- TorchVision
- Anaconda
- 실습
 - Installation
 - Tensors
 - Dataloader
 - Network
 - Training
 - Save and Load Model Weights
 - Transfer learning

Most of the slides are inspired by POSTECH
AIGS549 class (Prof. Suha Kwak)

PyTorch

- GPU와 CPU를 사용해서 deep learning 을 하기 위해 최적화된 tensor library이다.
- 원래 Lua라는 프로그래밍 언어로 된 Torch였으나, 오픈소스로 여러 사람이 개발하다가 facebook AI 가 개발에 참여하면서, 사용성이 좋아지고 있다.
- PyTorch로 넘어오면서 특히 좋은 점은 자동 미분 모듈이 있어서 따로 backward함수를 구현하지 않아도 된다는 점이다.
- TensorFlow 와 비슷하게 많이 사용되나, 본 강의에서는 PyTorch위주로 다룬다.
- 모듈, tensorflow와의 비교, 장점: <https://ko.wikipedia.org/wiki/PyTorch>
- 설치: <https://pytorch.org/get-started/locally/> : 각자 머신 스펙에 맞게 설치.
- PyTorch documentation: <https://pytorch.org/docs/stable/index.html>
- PyTorch tutorial: <https://pytorch.org/tutorials/>

TorchVision

Preliminary

Linear Algebra
Basic

- PyTorch project 의 일부로서 이미지를 처리하는 다양한 인터페이스(함수, 클래스 ...)를 제공한다.
- Numpy array (또는 PIL) 이미지의 torch.Tensor type으로의 형변환해주는 함수 (transforms.ToPILImage(), transforms.ToTensor())와,
- image를 자르거나 photometric jittering을 주는 등의 data augmentation 기능을 제공한다.
- 또한, 대표적인 CNN model과 그 pre-trained weights을 쉽게 사용하도록 해준다.
- 설치: 보통 pytorch 설치할 때에 같이 설치한다.
- TorchVision documentation: <https://pytorch.org/vision/stable/index.html>
- 추가적으로 image augmentation 할 때에 사용하는 albumentations 라이브러리가 있다.
- 예시: <https://albumentations.ai/>. Documentation: <https://albumentations.ai/docs/>

Anaconda (시작하기에 앞서)

- 아나콘다 : 패키지 (라이브러리 및 dependencies) 관리를 쉽게하기 위해 사용하는 오픈소스 패키지 관리 시스템.
- 도커나 다른 방식을 사용하기도 하지만, 보통 anaconda 를 사용한 패키지 관리가 간단하고 편리하기에 많이 사용한다.
- Conda 와 같은 package manager를 쓰면 프로젝트마다 패키지나 버전 관리가 간단하고, 프로그래밍 그 자체에 좀 더 집중할 수 있게 도와준다.
- 설치: <https://www.anaconda.com/products/individual>
- 위에서 installer 설치 후
bash Anaconda-latest-Linux-x86_64.sh 또는 GUI로 설치.
- 새로운 environment 만들기.
`conda create -n {desired name} python={desired version}`

Anaconda (환경 activate and remove)

- Using an environment
`source activate {desired env name}` `source deactivate`
- Remove an environment
`conda remove -n {desired env name} --all`
- More informations -> <https://docs.conda.io/projects/conda/en/latest/user-guide/install/index.html>

PyTorch and Torchvision 실습

Preliminary

Linear Algebra
Basic

- Installation
- Tensors
- Loading data
- Build the neural network
- Training / testing
- Model load and save
- Transfer learning (with torchvision model)

PyTorch 설치

- Installation according to your system environment



- Example: install pytorch on “Linux”, “Conda” with “CUDA 10.2”
conda install pytorch torchvision torchaudio cudatoolkit=10.2 -c pytorch
- *When you use conda, CUDA environment also be installed (check GPU compatibility).

Tensors: creation

- What is tensors
 - Core data structure in pytorch similar to arrays (Inputs, outputs, the model parameters).
 - Very similar to numpy ndarrays (but tensors can run on GPUs)
 - *If you're familiar with ndarrays, you'll already familiar with most of the Tensor API.
- Rest of the tutorial includes:
 - Tensor creation
 - Directly from data

```
import torch
import numpy as np
```

```
data = [[1, 2], [3, 4]]
x_data = torch.tensor(data) # 2x2 tensor
```

```
x_data = torch.tensor(3.14159) # scalar tensor
```

```
x_data = torch.tensor([]) # empty tensor
```

Tensors: creation

- From numpy arrays

```
np_array = np.array(data)
x_np = torch.from_numpy(np_array)
```

- Tensor creation using APIs

```
>>> shape = (2, 3,)
>>> torch.rand(shape)
tensor([[0.9864, 0.2620, 0.1767],
        [0.0322, 0.4785, 0.9585]])
>>> torch.ones(shape)
tensor([[1., 1., 1.],
        [1., 1., 1.]])
>>> torch.zeros(shape)
tensor([[0., 0., 0.],
        [0., 0., 0.]])
```

- Tensor creation from other tensors

```
x_ones = torch.ones_like(x_data) # retains the properties of x_data
x_ones = torch.rand_like(x_data, dtype=torch.float) # overrides the datatype of x_data
```

Tensors: attributes and GPU

- Tensor attributes (type이나 현재 처리되는 Device)

```
>>> tensor = torch.rand(3, 4)
>>> tensor.shape
torch.Size([3, 4])
>>> tensor.dtype
torch.float32
>>> tensor.device
cpu
```

- Run tensor on the GPU

```
>>> tensor = tensor.to('cuda')
>>> tensor.device
cuda:0
```

- ✓ operation runs on GPUs if the input tensor is on the cuda device
- ✓ all operand tensors to be on the same device

Tensors: operations

- Tensor operations

```
>>> tensor = torch.ones(4, 4)
>>> tensor[:,1] = 0 # indexing
>>> tensor
tensor([[1., 0., 1., 1.],
        [1., 0., 1., 1.],
        [1., 0., 1., 1.],
        [1., 0., 1., 1.]])
>>> torch.cat([tensor, tensor, tensor], dim=1) # joining
tensor([[1., 0., 1., 1., 0., 1., 1., 0., 1., 1.],
        [1., 0., 1., 1., 0., 1., 1., 0., 1., 1.],
        [1., 0., 1., 1., 0., 1., 1., 0., 1., 1.],
        [1., 0., 1., 1., 0., 1., 1., 0., 1., 1.]])
>>> tensor + tensor * tensor - tensor # element-wise arithmetic operations
tensor([[1., 0., 1., 1.],
        [1., 0., 1., 1.],
        [1., 0., 1., 1.],
        [1., 0., 1., 1.]])
>>> tensor @ tensor # matrix multiplication
tensor([[3., 3., 3., 3.],
        [3., 3., 3., 3.],
        [3., 3., 3., 3.],
        [3., 3., 3., 3.]])
```

Loading data

Datasets and DataLoader

- All datasets are subclass of `torch.utils.data.Dataset`
 - ✓ You can make your own datasets by inheriting `Dataset` class
 - ✓ Some famous datasets is provided in `torchvision.datasets`
- `DataLoader` reads datasets and handles multi-process mini-batch data loading

```
# Training cycle
for epoch in range(training_epochs):
    # Loop over all batches
    for i in range(total_batch):
```

▲ In the neural network terminology:

288

- one **epoch** = one forward pass and one backward pass of *all* the training examples
- **batch size** = the number of training examples in one forward/backward pass. The higher the batch size, the more memory space you'll need.
- number of **iterations** = number of passes, each pass using [batch size] number of examples. To be clear, one pass = one forward pass + one backward pass (we do not count the forward pass and backward pass as two different passes).

Example: if you have 1000 training examples, and your batch size is 500, then it will take 2 iterations to complete 1 epoch.

Loading data

Datasets from torchvision

- Example FashionMNIST

```
from torch.utils.data import Dataset
from torchvision import datasets
from torchvision.transforms import ToTensor
import matplotlib.pyplot as plt

training_data = datasets.FashionMNIST(
    root="data",
    train=True,
    download=True,
    transform=ToTensor())

test_data = datasets.FashionMNIST(
    root="data",
    train=False,
    download=True,
    transform=ToTensor())
```

- Caltech
- CelebA
- CIFAR
- Cityscapes
- COCO
- EMNIST
- FakeData
- Fashion-MNIST
- Flickr
- HMDB51
- ImageNet
- Kinetics-400
- KITTI
- KMNIST
- LSUN
- MNIST
- Omniglot
- PhotoTour
- Places365
- QMNIST
- SBD
- SBU
- SEMEION
- STL10
- SVHN
- UCF101
- USPS
- VOC
- WIDERFace

- The dataset will be downloaded by torchvision.
- More data augmentation techniques are in the reference page.

Find your dataset from

<https://pytorch.org/vision/stable/datasets.html>

https://pytorch.org/tutorials/beginner/data_loading_tutorial.html

Loading data

Custom dataset

- Methods you should define when inherit from `torch.utils.Dataset`

✓ `__init__(self, [])` : 데이터셋 루트 딕렉토리 / transform / 이미지 path ...

✓ `__len__(self)` : 데이터셋의 length (크기) 반환.

✓ `__getitem__(self, idx)` : (idx)th data 를 return. *Dataloader call this method

```
class DiabetesDataset(Dataset):
    """ Diabetes dataset."""

    # Initialize your data, download, etc.
    def __init__(self):
        1 download, read data, etc.

    def __getitem__(self, index):
        return
    2 return one item on the index

    def __len__(self):
        return
    3 return the data length

dataset = DiabetesDataset()
train_loader = DataLoader(dataset=dataset,
                        batch_size=32,
                        shuffle=True,
                        num_workers=2)
```

[ner/data_loading_tutorial.html](#)

Loading data

Dataloader

- Methods reads datasets and handles **multi-process mini-batch data loading**
 - ✓ Multi-process mini-batch loader can avoid bottleneck on reading and transferring data.
 - ✓ DataLoader is an iterable that abstracts this complexity for us in an easy API.

```
from torch.utils.data import DataLoader

train_dataloader = DataLoader(training_data, batch_size=64, shuffle=True)
test_dataloader = DataLoader(test_data, batch_size=64, shuffle=True)
```



<Needs to transfer data from storage into GPU>

Build the Neural Network

You should define two method in network class: “`__init__`” and “`forward`”

- - “`__init__`” : you need to implement the network architecture using `torch.nn` namespace .
- - “`forward`” : define how do you compute the forward path of the neural network.

```
class NeuralNetwork(nn.Module):
    def __init__(self):
        super(NeuralNetwork, self).__init__()
        self.flatten = nn.Flatten()
        self.linear_relu_stack = nn.Sequential(
            nn.Linear(28*28, 512),
            nn.ReLU(),
            nn.Linear(512, 512),
            nn.ReLU(),
            nn.Linear(512, 10),
        )

    def forward(self, x):
        x = self.flatten(x)
        logits = self.linear_relu_stack(x)
        return logits
```

Training

- Training data로 network 모델 학습하기.

1. Define Model
2. Define Dataloader
3. Define Optimizer
4. Iterate the below lines until reaching the max_epoch.
 - (1) Sample a batch from the data loader
 - (2) Predict the output using the neural network
 - (3) Compute the error between the answer and prediction
 - (4) Compute the gradient using optimizer.
 - (5) Backpropagate the gradient to optimize the NN parameters.

Training

1. Define Model
2. Define Dataloader
3. Define Optimizer
4. Iterate the below lines until reaching the max_epoch.
 - (1) Sample a batch from the data loader
 - (2) Predict the output using the neural network
 - (3) Compute the error between the answer and prediction
 - (4) Compute the gradient using optimizer.
 - (5) Backpropagate the gradient to optimize the NN parameters.

Preliminary

Linear Algebra
Basic

```
def train_loop(dataloader, model, loss_fn, optimizer):  
    size = len(dataloader.dataset)  
    for batch, (X, y) in enumerate(dataloader):  
        # Compute prediction and loss  
        pred = model(X)  
        loss = loss_fn(pred, y)  
  
        # Backpropagation  
        optimizer.zero_grad()  
        loss.backward()  
        optimizer.step()  
  
        if batch % 100 == 0:  
            loss, current = loss.item(), batch * len(X)  
            print(f"loss: {loss:>7f} [{current:>5d}/{size:>5d}]")
```

세 줄로 이루어지는
optimization step.

Save and Load Model Weights

- Pytorch library also provide the functions(**torch.save**, **torch.load**) to save and store model weights.

```
model = models.vgg16(pretrained=True)
torch.save(model.state_dict(), 'model_weights.pth')
```

```
model = models.vgg16() # we do not specify pretrained=True, i.e. do not load default
weights
model.load_state_dict(torch.load('model_weights.pth'))
model.eval()
```

Transfer learning (with torchvision model)

- 모델을 scratch 부터 학습하는 것은 데이터도 많이 필요하고, 학습이 수렴하기까지 시간도 오래 걸린다.
- 이 때, pre-trained model을 사용하면 효율적으로 모델 학습을 할수 있다.
- 다음은 pre-trained model 을 사용한 image classification model fine-tuning (transfer learning) 예시이다.

```
model_ft = models.resnet18(pretrained=True)
num_ftrs = model_ft.fc.in_features
# Here the size of each output sample is set to 2.
# Alternatively, it can be generalized to nn.Linear(num_ftrs, len(class_names)).
model_ft.fc = nn.Linear(num_ftrs, 2)

model_ft = model_ft.to(device)

criterion = nn.CrossEntropyLoss()

# Observe that all parameters are being optimized
optimizer_ft = optim.SGD(model_ft.parameters(), lr=0.001, momentum=0.9)

# Decay LR by a factor of 0.1 every 7 epochs
exp_lr_scheduler = lr_scheduler.StepLR(optimizer_ft, step_size=7, gamma=0.1)
```

컴퓨터 비전 초격차: 이론/실습

6 Preliminary – logging, wandb, kornia

6

Index

- logging
- wandb
- kornia

logging

- 컴퓨터비전 / 딥러닝 모델 개발 뿐 아니라 모든 프로그래밍에 있어서 logging 모듈을 잘 구현하는 것은 매우 중요하다.
- Python 에도 logging library 가 있다.
- 다양한 이벤트 수준에 대해서 정의할 수 있고, logger 와 handler가 이를 적절히 처리해 준다.
- 로깅에 대해 더 자세히 설명하는 것은 이 강의의 범위를 벗어나므로 생략한다.
- 설치: pip install logging
- Documentation: <https://docs.python.org/ko/3/howto/logging.html>

wandb

- Wandb (weights & Biases)는 학습 모델 개발 시에 하이퍼파라미터나, 시스템 메트릭, 결과를 로깅해주는 패키지이다.
- 비슷하게는 tensorboard/tensorboardX나 hydra 같은 패키지들이 있다.
- Wandb의 장점은 웹서비스를 통해서 다양한 모델의 결과를 확인/관리 할 수 있다는 점이다.
- 또한, sweep이라는 기능은 autoML (grid search, Bayesian optimization 등)을 통해 하이퍼파라미터 서치를 지원해준다. 베스트 모델을 찾기 위해서 하이퍼파라미터 서치가 필수적인데, 이때 매우 유용하다.
- 설치 : pip install wandb
- Documentation: <https://docs.wandb.ai/>

Kornia

- Kornia 는 AI를 위한 컴퓨터비전 알고리즘의 State-of-the-art 구현체를 모아놓은 라이브러리로 비교적 최신인 2019년부터 개발되기 시작하였다.
- 특히 딥러닝 기반의 컴퓨터비전 모델을 연구/개발할 때 유용하다.
- 구현하기에 까다로운 보조함수들이 구현되어 있고, GPU 기반의 deep learning 프레임워크인 PyTorch와도 호환가능하다.
- 설치: pip install kornia
- Documentation: <https://kornia.readthedocs.io/en/latest/>
- Github: <https://github.com/kornia/kornia>

